

---

# **Distributed Coordination and Agreement**

---


# Coordination and Agreement

---

- What is **coordination**?
  - Coordinate the activities/decision of different processes (process **synchronization**)
- Coordination among two types of processes
  - Coordinating processes belonging to the same parent process
    - Distributed processes may need to coordinate with each other to complete a task
    - i.e., When to start and when to end
  - Independent processes from different applications
    - Coordination in accessing common resources, i.e., global data (data synchronization)
- What is an **agreement** (**consensus**)
  - All processes make the same decision
  - i.e., all processes /threads agree to commit or abort


# Mutual Exclusion and Critical Section

- Coordination is required to access shared resources to ensure data consistency (data synchronization)
  - Mutual exclusion prevents inconsistency among concurrent processes
- Mutual exclusion requirements:
  - Only one process is allowed to access the shared resource at a time
  - When a process is using a resource (i.e., shared global variables), other processes requesting the resource have to wait
- How to achieve mutual exclusion?
  - Mutual exclusion can be achieved by defining critical sections



# Distributed Mutual Exclusion

- Distributed mutual exclusion problem:
  - Multiple processes in different locations need to access a shared resource
  - A global resource may be replicated at multiple locations and managed by multiple servers
  - Distributed critical sections



The deadlock problem

# Operations for Critical Sections

---

- Operations for accessing a shared resource:
  - Enter(): enter a critical section; otherwise it is blocked
  - ResourceAccess(): access shared resources in critical section
  - Exit(): leave a critical section

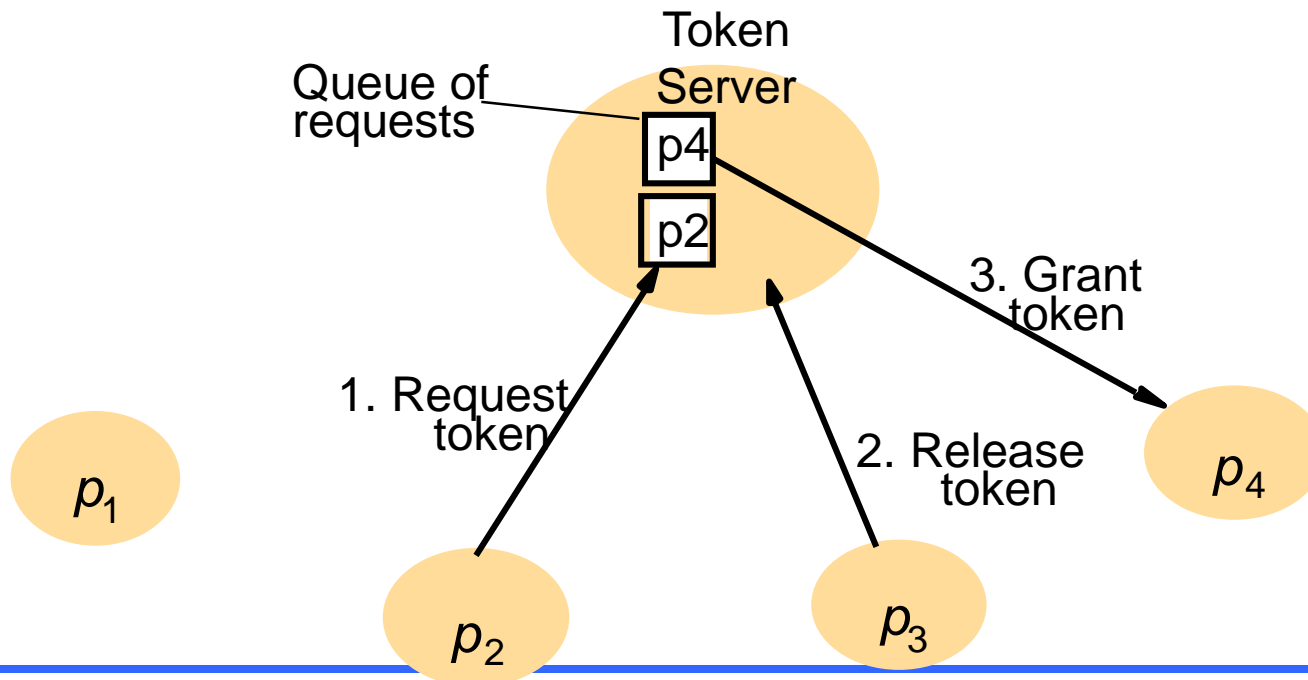
# Algorithms for Mutual Exclusion

---

- Performance metrics
  - **Communication cost** (number of messages generated)
  - **Access delay**: waiting time to access to a resource
  - **Synchronization time**: minimum time from the release of a resource to the next assignment of the resource (someone is waiting)
- Other performance concerns
  - Deadlock
  - Starvation: infinite postponement (waiting forever)
  - Fairness: fair for all processes waiting for entering a critical section (i.e., following the arrival order of the requests)


# A Central Server Algorithm

- Algorithm with a Central Token-server
  - A centralized server is responsible for granting token to access the resource
  - To access to a shared resource, send a request to the server with the token
  - Note: **The server only manages the token, but not the resource (2-in-1 is ok)**



# Performance of Central Server Algorithm

- What is the queuing/scheduling method for waiting requests?
  - FIFO, priority-based, etc.
- What is the performance of the method?
  - Entering the critical section costs: 2 messages
  - Exiting the critical section needs: 1 message
  - The minimum delay for access to the resource is one round-trip delay
  - The synchronization delay is also a round-trip delay:
    - a release message to the server followed by a grant message
- Other issues:
  - The central server becomes the bottleneck
  - Not scalable






# A Distributed Method: Token-Ring Algorithm

- All processes are organized in a logical ring. A token is circulated from one process to the next along the ring.
- Only the process that holds the token is allowed to access the resource.
  - The process passes the token to the next once it finishes the access.
- A process will pass the token immediately to the next if does not need to access the resource.

**Assume: no process or network failure**



# Performance of Token-Ring Algorithm

---

- Performance issues:
  - It costs network messages constantly (regardless whether there is any process requesting the resource)
  - The access delay is from 0 message to N messages depending on the location of the token at the time when it generates the request
    - Best case: when the token is passed to P1 and P1 just wants to access the resource (0 waiting)
    - Worst case: right after P1 releases the token, P1 wants to access to the resource (N-hops waiting)
  - The synchronization delay is also from 0 message to N messages

# Ricart and Agrawala's Algorithm: Using Multicast and Timestamp

---

- When a process wishes to access a resource, it multicasts a request message to all other processes in the group
  - Assumption: network is interconnected and the communication is reliable
- Any process that receives the request replies “ok”
- The process can access the resource only when it receives replies from **all other processes**
- No two processes will receive permissions from all others at the same time (mutual exclusion is guaranteed)
- **How about two processes multicast requests simultaneously?**

# Resolve Concurrent Requests using Timestamps

- When a process multicasts a request, it attaches its own ID and its current timestamp  $\langle T_i, p_i \rangle$  to the message
- In case there are multiple concurrent requests, when a process receives a request:
  - It replies a request only if itself is not waiting to access OR itself is waiting **but** the timestamp of the request is smaller than its own request-time
    - Earlier request has a higher priority (FCFS)
  - All requests are serialized according to their timestamps
  - **Assume all clocks of processes are strictly synchronized**

# Algorithm of using Multicast and Timestamp

- A process is in one of the following states
  - RELEASED: outside the critical section
  - HOLD: inside the critical section (hold the resource)
  - WAIT: waiting to enter the critical section
- When  $p_i$  wishes to enter a critical section:
  - It sets its status to WAIT and multicasts  $\langle T_i, p_i \rangle$  to all other processes
  - It blocks until receiving all replies before entering the critical section
- Upon receiving a request message, if this process is in:
  - RELEASED, it replies to  $p_i$  immediately
  - HOLD, it delays the reply to  $p_i$  until it exits from the critical section
  - WAIT, it compares its timestamp with the one in the message. If its own timestamp is greater, it replies to  $p_i$  immediately; otherwise, it delays the reply until it exits from the critical section

# Ricart and Agrawala's Algorithm: Using Multicast and Timestamp

## *On initialization*

$state := \text{RELEASED};$

## *On wishing to enter the critical section*

$state := \text{WAIT}; T := \text{request's timestamp};$

Multicast *request* to all processes;

## *On receiving a reply from $p_i$*

if the number of replies received  $== (N - 1)$

$state := \text{HOLD};$

enter the critical section (access the resource);

end if

## *On $p_j$ receiving a request $\langle T_i, p_i \rangle$ from $p_i$ ( $i \neq j$ )*

if ( $state = \text{HOLD}$  or ( $state = \text{WAIT}$  and ( $T_j < T_i$ )))

queue *request* from  $p_i$  without replying;

else

reply  $p_i$  immediately;

end if

## *On exiting the critical section*

$state := \text{RELEASED};$

reply to all queued requests; // unblock all others

# Performance of Ricart and Agrawala's Algorithm

---

- Performance issues:
  - Gaining entry into the critical section requires  $2(N - 1)$  messages
  - Minimum delay in accessing: one round-trip time ( $2d$ )
  - Synchronization delay:  $d$  (once a process leaves a critical section, it replies to all waiting processes)
  - ? Deadlock
  - ? Starvation

# Motivation of Maekawa's Voting Algorithm

- Problems of the multicast algorithm
  - Large number of synchronization messages (1 multicast and N replies)
  - Have to get the permission from all the member processes
  - To enter a critical section, there is no need to get permission from all peers
- Improvement by Maekawa's Algorithm
  - A process only needs to obtain permissions from a subset of its peers (votes), so long as **any two subsets always have overlap**
  - **Why?** A process can give permission to only one process. The overlapping member would prevent two processes from entering at the same time
- How to determine the size of the subset (called *quorum*):
  - Any two subsets must have at least one common member (process)
  - A process **must** obtain sufficient votes (quorum) to enter the critical section



# Prof. Maekawa



# Theory of Maekawa's Voting Algorithm

- Each  $p_i$  is associated with a voting set  $V_i$  ( $p_i$  needs to obtain permissions from all processes in  $V_i$ ),  $p_i \in V_i$
- $V_i \cap V_j \neq \Phi$  for all  $i, j = 1, 2, \dots, N$ 
  - The overlapping element prevents two processes from entering a critical section at the same time
  - e.g.,  $V_1: \{p_1, p_2\}$  and  $V_3: \{p_2, p_3\}$ .  $p_2$  will not grant permission to  $p_3$  if it has already voted for  $p_1$
  - Majority voting (an easy example)
- $|V_i| = K$ : all processes require the same number of votes (quorum)
  - Minimizing  $K$  can improve the performance as the number of messages for synchronization is reduced
  - The minimal  $K$  for mutual exclusion:  $K \sim \sqrt{N}$  – why?

# Maekawa's Voting Algorithm

- When process  $p_i$  wishes to enter a critical section:
  - $p_i$  multicasts a *request* message to all members in  $V_i$  (including itself)
  - $p_i$  is blocked until it receives all replies from the members in  $V_i$
- When  $p_j$  in  $V_i$  receives  $p_i$ 's request: if it is **HOLD** or it has already replied to (voted for) another process (including itself), it queues  $p_i$ 's request; otherwise it replies  $p_i$  immediately
- When  $p_i$  exits the critical section, it sends a *release* message to all members in  $V_i$
- When  $p_j$  receives a *release* message, it removes the head of the queued requests and replies to it
  - Serving the requests one by one

# Maekawa's Algorithm

## *On initialization*

*state* := RELEASED;

*voted* := FALSE;

## *For $p_i$ to enter the critical section*

*state* := WAIT;

*voted* := TRUE; // vote for itself

Multicast *request* to all processes in  $V_i$ ;

Wait until (No. of replies received =  $K$ );

*state* := HOLD;

## *On receipt of a request from $p_i$ at $p_j$*

if (*state* = HOLD or *voted* = TRUE)

    queue *request* from  $p_i$  (no reply);

else

    send *reply* to  $p_i$ ;

*voted* := TRUE;

end if

## *For $p_i$ to exit the critical section*

*state* := RELEASED;

Multicast *release* to all processes in  $V_i$ ;

## *On receipt of a release from $p_i$ at $p_j$*

if (queue of requests is non-empty)

    remove head of queue, say  $p_k$ ;

    send *reply* to  $p_k$ ;

*voted* := TRUE;

else

*voted* := FALSE;

end if

# An Example of Maekawa's Voting Algorithm

---

V1: {P1, P2, P4}

V2: {P2, P3, P4}

V3: {P1, P2, P3}

V4: {P2, P3, P4}

- 1) If P1 wishes to enter a critical section, it multicasts a *request* to itself, P2 and P4. When it receives replies from all of them, it enters the critical section
- 2) In the meantime, if P3 wishes to enter, it multicasts a *request* to itself, P1 and P2. It will be blocked by P1 and P2.

# Maekawa's Voting Algorithm

---

- Performance Issues
  - $O(\sqrt{N})$  messages per entry into the critical section – why?
  - $O(\sqrt{N})$  messages per exit from a critical section
  - Synchronization delay is round-trip ( $2d$ )
- The algorithm is deadlock prone - why?
  - How to solve it?

# King's Poisoned Wine Problem

- A King prepared 100 barrels of wine to host a big banquet with his international guests
- Before the banquet starts, the King got the information that **one barrel of the wine was poisoned**
- Design a method to use the minimum number of prisoners to sample the wine and find out the poisoned barrel of wine



# Election Algorithms

---

- Many distributed applications require one process act as a leader, such as the central server algorithm for mutual exclusion. How to elect the leader?
- Assume that each process has a unique ID (used as the priority). **The process with the highest ID will be elected**
- The goal of election algorithms: when election completes, all **participating** processes agree on who the new leader is
- Any process can initiate the election at any time
  - e.g., when a process detects the failure of the current leader, or feels there is a need to elect a new leader



# Requirements of Election Algorithms

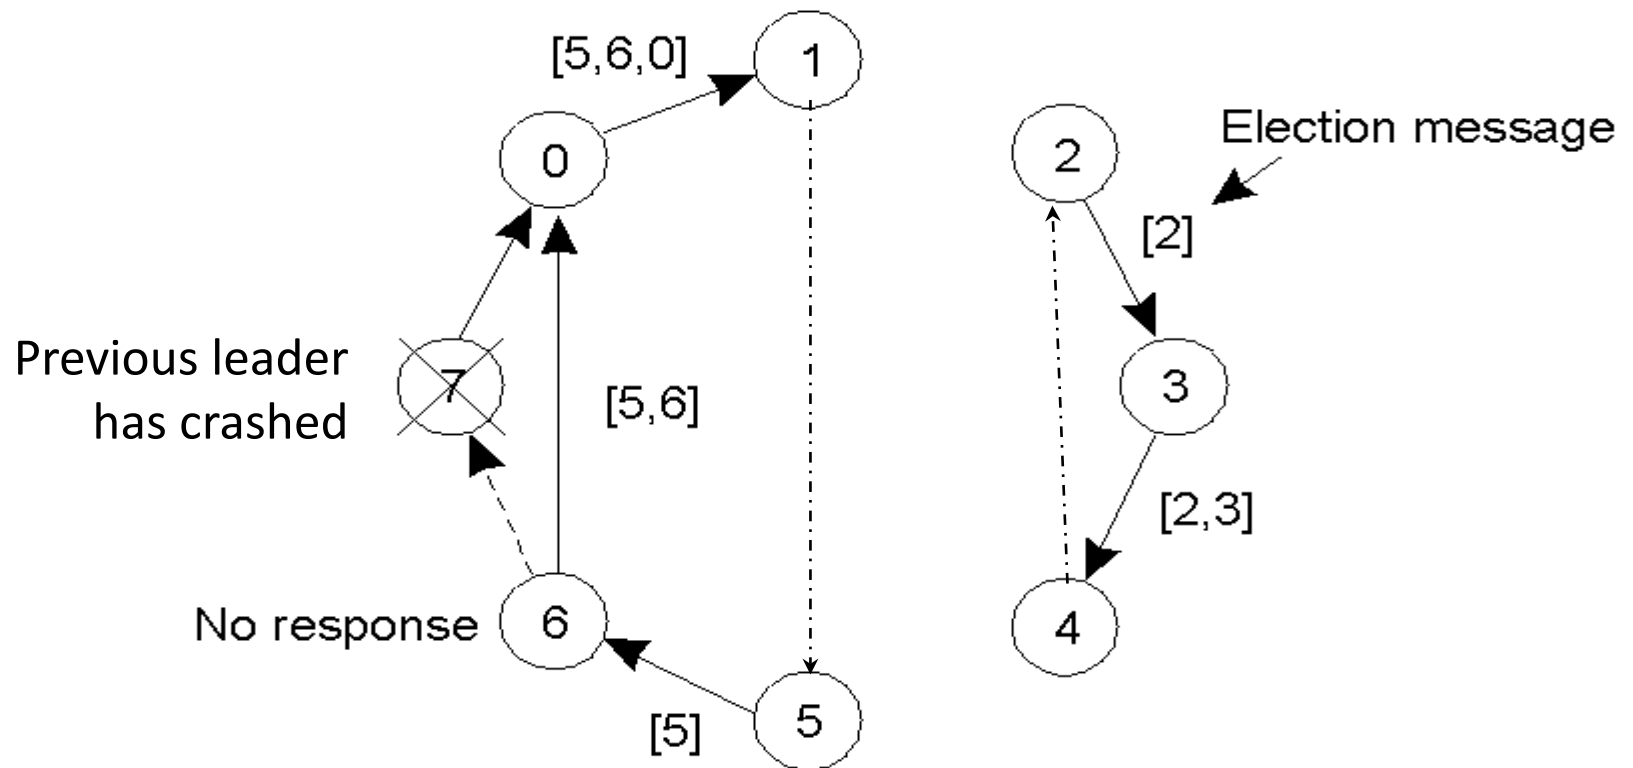
- Each process  $p_i$  ( $i = 1, 2, \dots, N$ ) has a variable  $electd_i$  that is the ID of the elected leader.
  - $electd_i$  is initially set to  $\Phi$  (*EMPTY*) when  $p_i$  participates in an election,  $i = 1, 2, \dots, N$
- Requirements
  - E1 (safety): **A participating process**  $p_i$  has either  $electd_i = \Phi$  or  $electd_i = P$  ( $P$  is the elected leader)
  - E2 (liveness): **All participating processes**, say  $p_i$ , eventually set  $electd_i \neq \Phi$  or crash
- Performance:
  - Number of messages for election
  - Turnaround time for election (No. of rounds of message exchanges)

# Election Algorithm: Ring-based Algorithm

- All processes are organized in a logical ring (similar to the token-ring)
- When a process notices that the leader is not functioning, it initiates the election by sending an *election* message containing its ID to its successor along the ring. **If the successor is down, skips over it** and goes to the next one, or the next after that, until a live process is located
- When a process receives an election message, if its own ID is greater than the one in the message:
  - 1) replaces the ID in the message by its own (**or adds its ID to the list**)
  - 2) forwards the message to its successor
- If a process receives the same election msg again and its ID is the greatest:
  - 1) sets its status to be the leader and  $elected_i \leftarrow$  its ID
  - 2) informs all processes by circulating  $elected_i$  message along the ring
- The election is complete when the  $elected_i$  msg reaches the original sender


# The Ring-based Algorithm: an example

Election algorithm using a ring with process failures



# Performance of Ring-based Algorithm

- Performance Issues
  - $N - 1$  messages for an initiator to reach the process with the highest ID in **worst case**
    - When the process with highest ID is next to the initiator in anti-clock wise (eg, P0), it takes  $N - 1$  msgs to reach this highest ID process
  - $N$  messages to circulate the highest ID
    - The process with highest ID knows it wins the election by now
  - $N$  messages to inform all members
    - For announcing the result
  - Totally  $3N - 1$  messages and delays (worst case)




# Election: The Bully Algorithm


Assume: every process knows the process IDs of others

- When process  $p_i$  discovers the failure of the leader, it initiates an election by sending an *election* message to all processes that have higher IDs than itself
- When process  $p_j$  receives the message from a lower-ID process, it replies “*I’m alive*” (OK) message to the sender and it takes over the election by sending out an election message to higher-ID processes (the same as  $p_i$ )
- If  $p_i$  receives no reply, it wins the election and becomes the leader:
  - The new leader announces its victory by sending all processes a message telling them that it is the new leader

# The Bully Algorithm: an example




(a)





Previous leader  
has crashed

(b)



(c)

## The Bully Algorithm: example (cont'd)

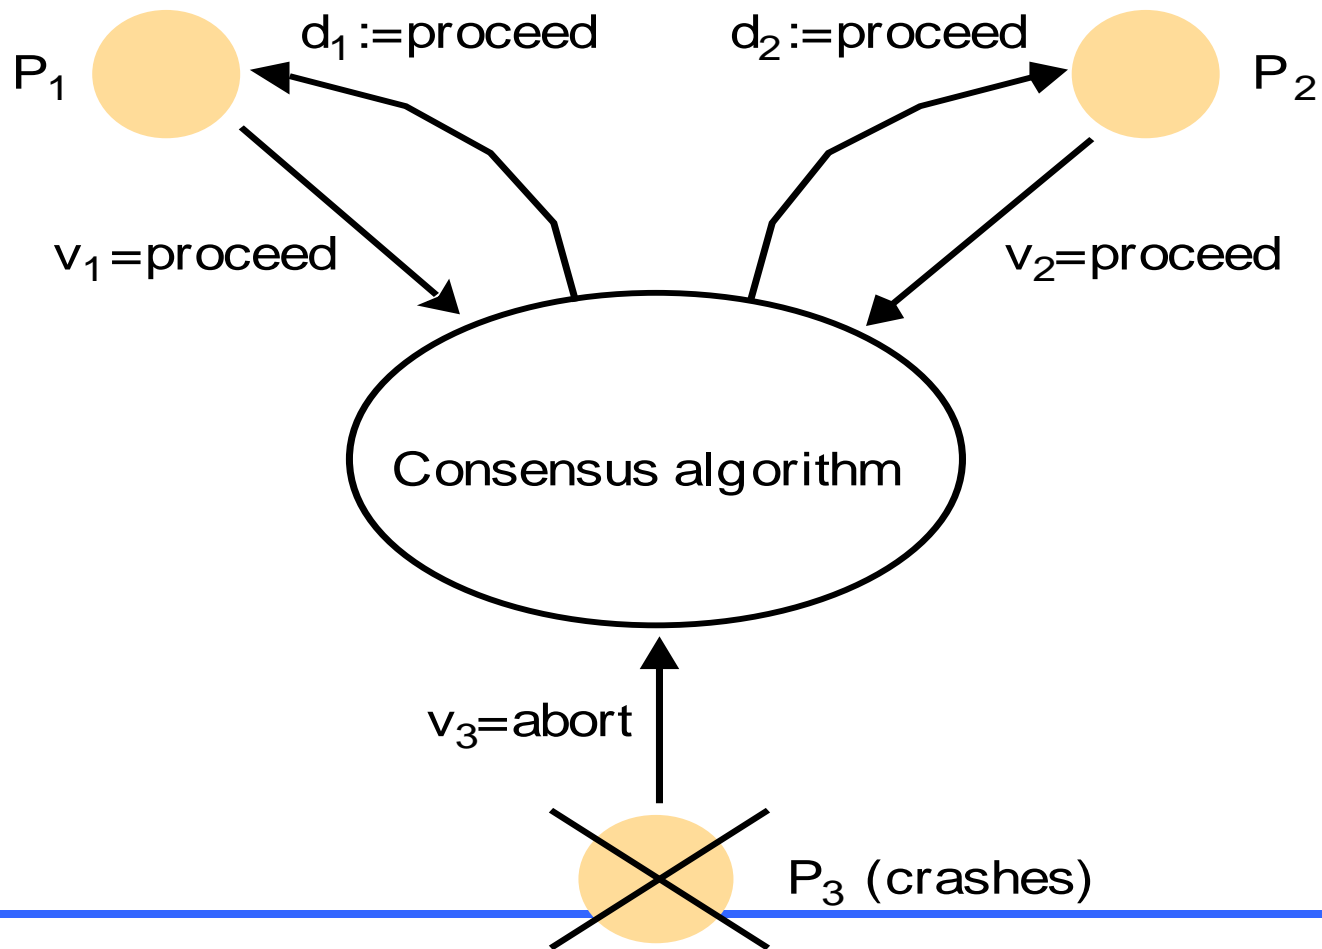


# Consensus Problem and Requirements

- Consensus: a group of processes agree on a value/decision after one or more of them propose the values/decisions
- Consensus problem:
  - Each process  $p_i$  begins with the *undecided* state and proposes a value  $v_i$  drawn from a set  $D$  ( $i = 1, 2, \dots, N$ )
  - All processes exchange their values with each other
  - Each process then decides the final value  $d_i$  based on the values proposed by others
- Requirements:
  - **Agreement**: all **correct processes** eventually have the same decision value
  - **Integrity**: if the correct processes all proposed the same value, any process in the decided state must choose this value



# Consensus of Three Processes



# A Simple Consensus Algorithm: **No process failure**

- 1) In a group of  $N$  processes, each process **multicasts** its proposed value to all other members in the group
  - 2) Each process waits until it has collected all  $N$  values including its own
  - 3) It then evaluates a function, i.e.,  $majority(v_1, v_2, \dots, v_n)$ ,  $max$  or  $min$ , *etc.*, to arrive a final value for the decision
- Termination is guaranteed by reliable communication
  - Agreement and integrity are guaranteed (all processes receive the same set of values and follow the same *majority* function)
  - But, if processes can fail, ...

# Consensus in Synchronous System: processes can (crash) fail

Assume: up to  $f$  out of  $N$  processes can suffer “crash” failure

- 1) Initial round: each process  $p_i$  includes its own value in  $V_i^1$  and multicasts  $V_i^1$  to all group members
- 2) The  $r^{\text{th}}$  round: each  $p_i$  collects values multicast from other group members, adds them into  $V_i^{r+1}$ , and multicasts  $V_i^{r+1} - V_i^r$ .
- 3) Repeat step 2) until  $r > f + 1$ . Each  $p_i$  returns  $d_i = \text{Min}\{V_i^{f+1}\}$

Note:

- “multicast” is not atomic: a process can fail in the middle of a multicast, leaving some receive but others not receive the msg
- At the end of  $f + 1$  round, every alive process receives the same set of values (some values could be proposed by failed processes)

# Algorithm of Consensus in Synchronous System of up to $f$ (crash) failures

Algorithm for process  $p_i \in G$ ;

*On initialization*

$Values_i^1 := \{v_i\}; Values_i^0 = \{\}$ ;

*In round  $r$  ( $1 \leq r \leq f + 1$ )*

$multicast(G, Values_i^r - Values_i^{r-1});$  //send only values that are not sent before

$Values_i^{r+1} := Values_i^r;$

*While* (in round  $r$ ) { // collect values multicast by other processes in  $r$ th round

*receive value*  $V_j$  from  $p_j$ ;

$Values_i^{r+1} := Values_i^{r+1} \cup V_j;$

}

*After  $(f + 1)$  rounds*



$Assignd_i = \min(Values_i^{f+1});$

# Consensus with Arbitrary Failures: The Byzantine Generals Problem

A command and generals exchange msgs to agree on attack or retreat

- The commander issues an order to his generals on attack or retreat
  - The commander could be faulty too
  - The generals exchange msgs among themselves about what they hear
- **Arbitrary failure model:** A faulty general (or the commander) can propose attack to some of the generals but retreat to others
- **Agreement:** The decision by **all correct processes** is the same
- **Integrity:** if the commander is correct, **all correct processes** decide on the value that the commander has proposed
- **Conclusion:** Byzantine generals problem has a solution iff  $f < N/3$ 
  - Faulty processes must be less than  $N/3$
  - Impossibility of 3 processes with 1 faulty

# The Byzantine Generals Problem: impossibility of 3 processes



Faulty processes are shown coloured

# Algorithm for Byzantine General Problem

## Assumptions:

- A1: every message is delivery correctly
- A2: the receiver of a message knows who sent it
- A3: the absence of a message can be detected

Definition. majority function:  $\text{majority}(v_1, \dots, v_{n-1})$  returns:

- 1) The majority value of among  $\{v_1, \dots, v_{n-1}\}$  if it exists; otherwise “retreat”
- 2) The median value of the ordered set  $\{v_1, \dots, v_{n-1}\}$

Note: the default value is “retreat” if a process doesn’t receive any value from another (or commander)

# The algorithm

Suppose 1 commander,  $n - 1$  generals and  $m$  traitors, algorithm  $\text{BGP}(m)$  is a recursive function:

## $\text{BGP}(0)$

- 1) The commander sends his value to every general
- 2) Each general uses the value he receives, or “retreat” as default

## $\text{BGP}(m)$ // $m$ is the number of traitors

- 1) The commander sends his value to every general
- 2) Each general  $p_i$ , receiving value  $v_i$  from commander, acts as the commander to call  $\text{BGT}(m - 1)$  sends  $v_i$  to all  $n - 2$  other generals
- 3) Each  $p_i$ , receiving value  $v_j$  from  $p_j$  in step 2 (using  $\text{BGT}(m - 1)$ ), uses value  $\text{majority}(v_1, \dots, v_{n-1})$



# Summary

---

- Distributed mutual exclusion
  - Central algorithm
  - Token-ring algorithm
  - Distributed Algorithm (Ricart and Agrawala's)
  - Voting Set Algorithm (Maekawa's)
- Elections
  - Ring-based Algorithm
  - Bully Algorithm
- Consensus
  - Distributed Consensus with crash failures
  - Byzantine Generals Problem

# Exercise

---

- 1a) Explain why no two processes can enter critical section at the same time with Maekawa's Algorithm.
- 1b) Maekawa's algorithm can reduce the communication cost to  $O(\sqrt{N})$  for a process to enter or exit from the critical section, where  $N$  is the total number of processes in the system. Why?
- 2) In the Consensus Algorithm that can tolerate at most  $f$  process failures (crashes), it requires  $(f+1)$  rounds of collecting & multicasting values before all correct processes reach consensus. Why?