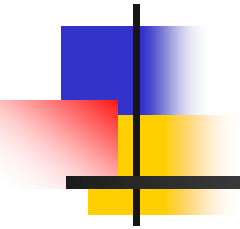

# BitCoin and Blockchain

# The rise of cryptocurrencies

- Bitcoin Price (USD) – Source : coinbase.com


- Bitcoin sparked research into multiple challenging areas and applications
  - more than 2000+ cryptocurrency startups according to angel.co

# Properties of Bitcoin

- Decentralization
  - no central authority that controls the entire network
- Non-repudiation
  - participants in the bitcoin network cannot deny their transactions
- Immutability
  - once a transaction is written into the ledger (i.e., blockchain), it cannot be altered
- Pseudonymous
  - no association between bitcoin participants and real-world identities

# A centralized ledger

- E-ledger (list of transactions)
  - a ledger in bitcoin is to trace the transaction-history of a coin
    - No balance of a person/account appears anywhere
  - but, a ledger in a bank maintains the current balance all the time

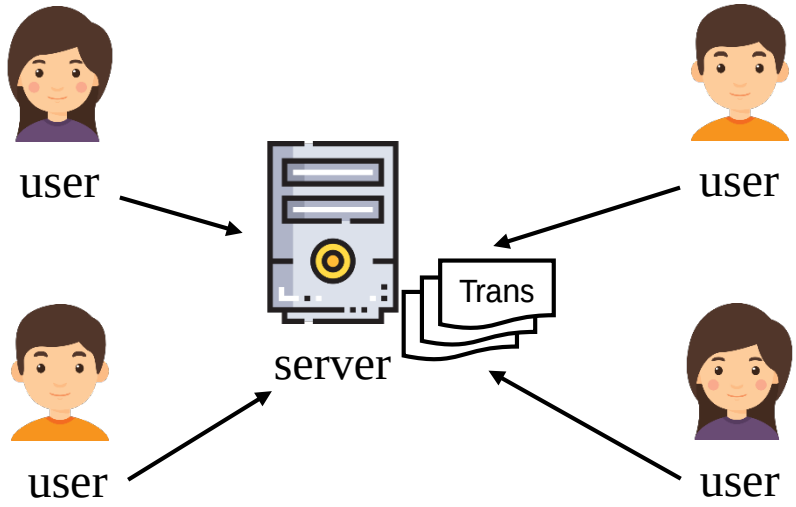


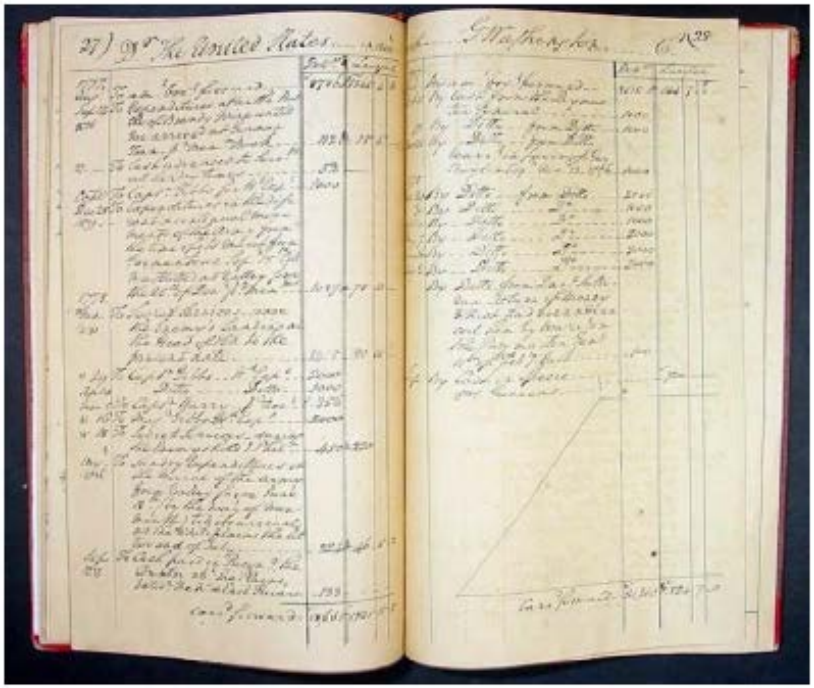

# Non-repudiation: sign transactions

- Digital coin == chain of digital signatures
- Ownership transfer of a coin:
  - Each person is identified by his public key in the cyber world
  - *A* transfers a coin to *B*: *A* signs the trans. using its private key
  - Sign(*Prev trans* + *New owner's public key*) // '+': concatenate 2 msgs
- Anyone can verify the transfer from the (n-1)$^{th}$ owner to the n$^{th}$

- But, who is responsible to keep the history of transactions in a decentralized system?

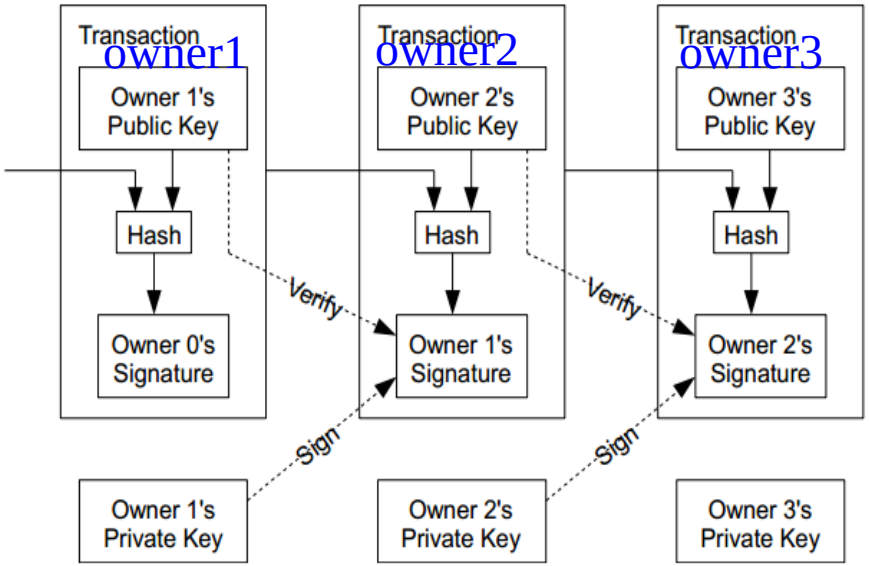

# A decentralized ledger

- Challenges in decentralized systems:
  - no authority keeps the transaction history
  - people may fake a transaction or double spend a coin by taking advantages of network delay
- To prevent fraud and double-spending:
  - each transaction is broadcast to all nodes
  - a transaction is confirmed only after verification (by whom?)

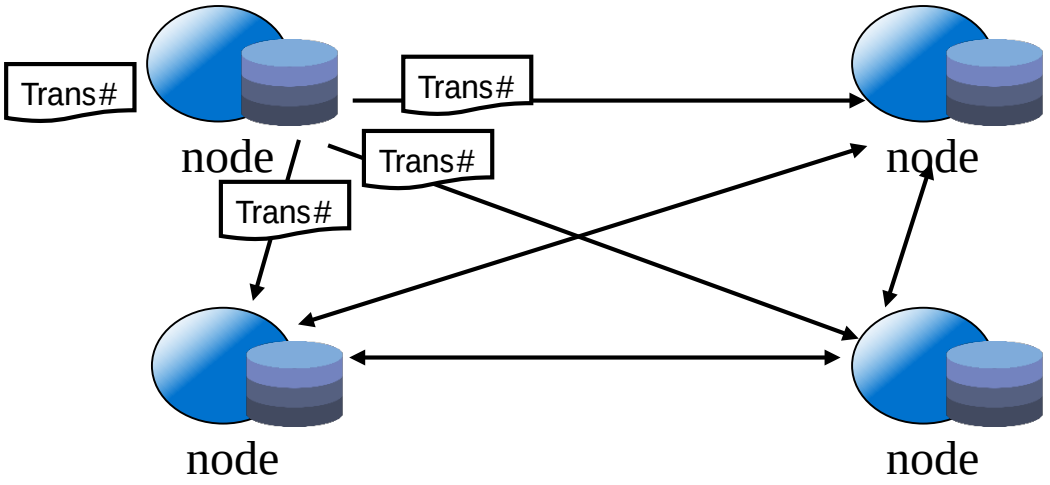

# Consistency of a distributed ledger

- Nodes receive different sets of trans at any time-point due to different network delays
- How to organize and verify the transactions to make a consistent distributed ledger?

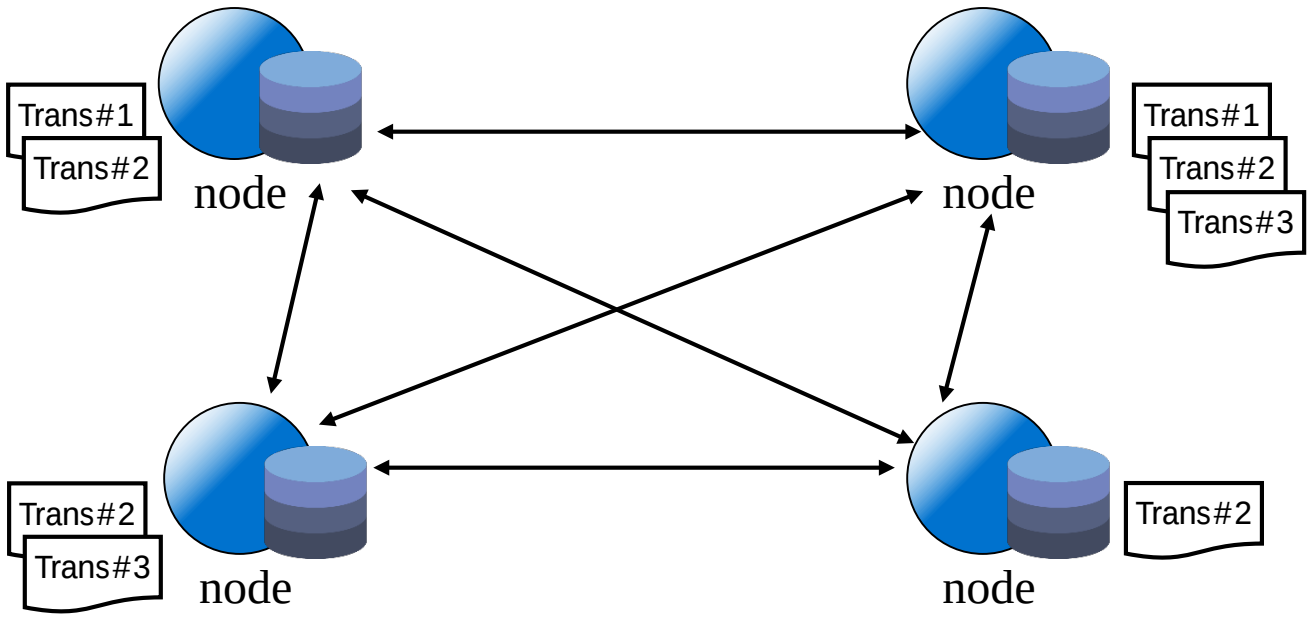

# Blockchain and transactions

- Blockchain, a chain of blocks, is a distributed ledger, recording all trans in the system
  - each block contains of a set of verified trans
- Each node (mining node) selects a set of trans from its local pool, verifies them, generates a new block, and links the new block to the chain
- Other nodes, upon receiving this new block, will accept the new block by further linking their new blocks to it
  - by "accept a block", it means to verify the trans again in the block to prevent the creator of the block from making any fraud trans

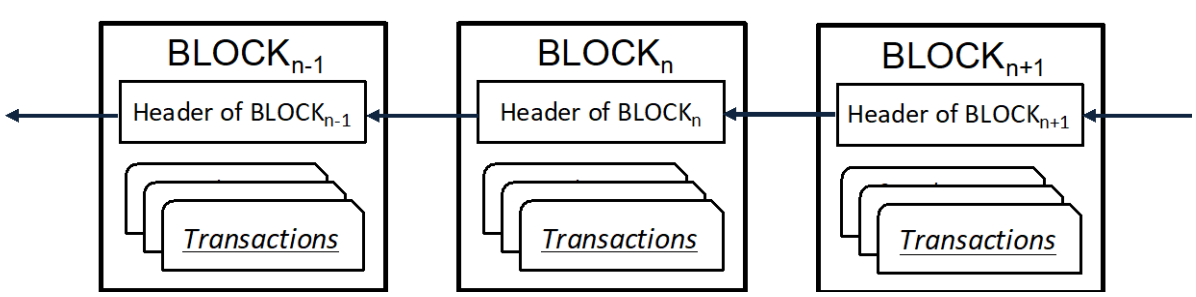


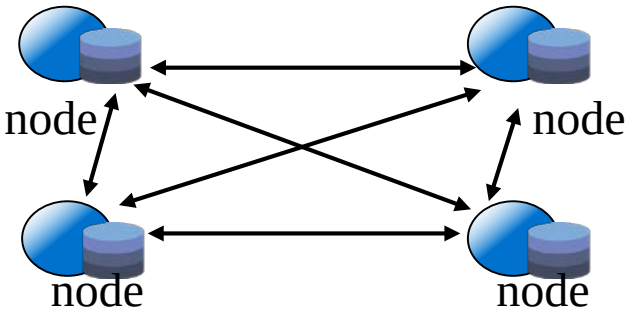

# Transaction flow

- The coin was created for Goofy by the *Coinbase transaction* (discussed later) and Goofy is the owner
- Transactions over a coin are chained up

A coin's owner can spend/transfer it

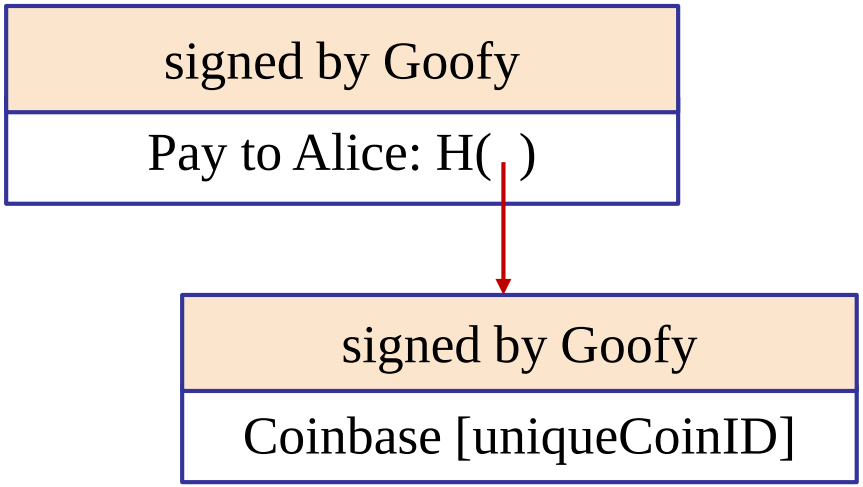

# Chain of transaction flow

The recipient can pass on the coin again

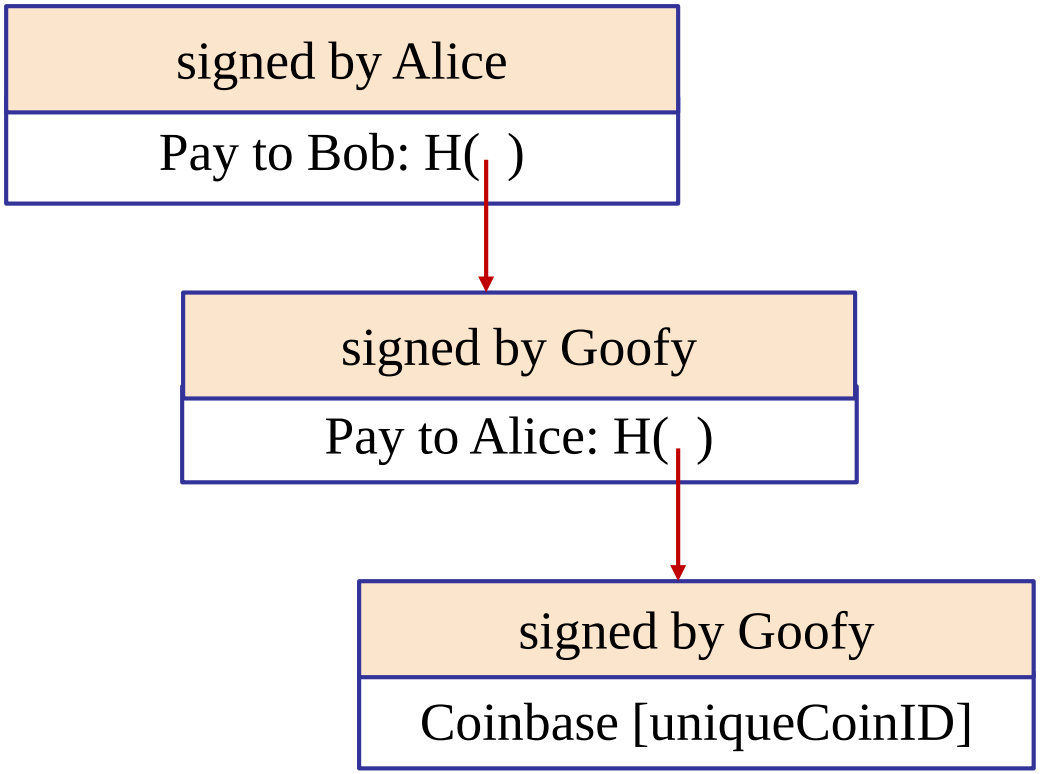

# Full transaction chain: A ledger

- The full chain is a complete ledger/ history of all trans
  - the input of the current trans points to the output of an earlier trans, indicating the source of the trans
- The history of the full blockchain reveals the state/ ownership of all bitcoins (BTC)
- The ledger is structured in terms of transactions
  - no explicit “account balance”

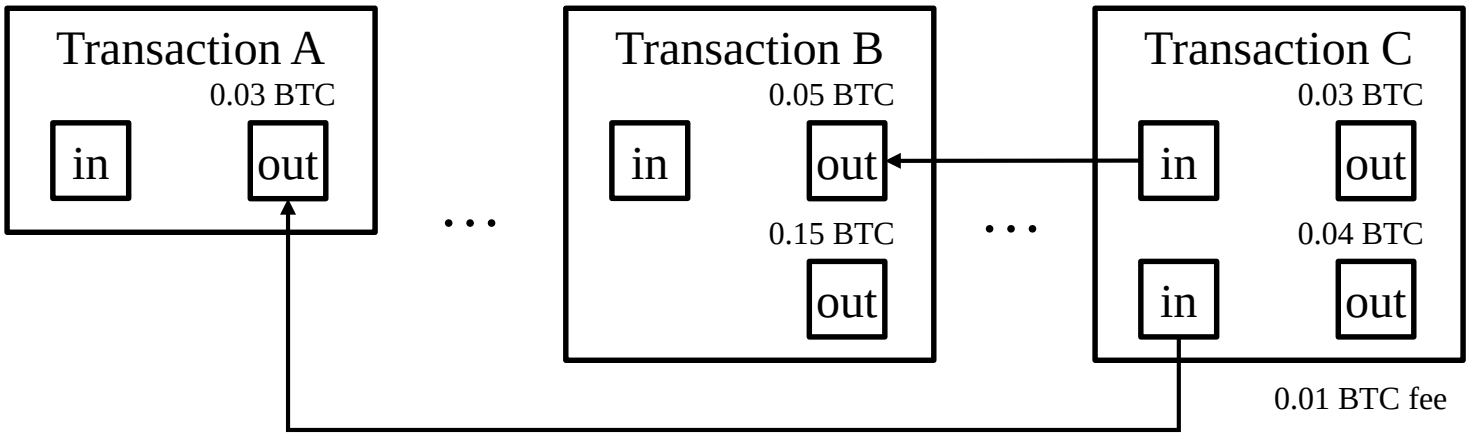

# Trans-based ledger: without in/out pointer

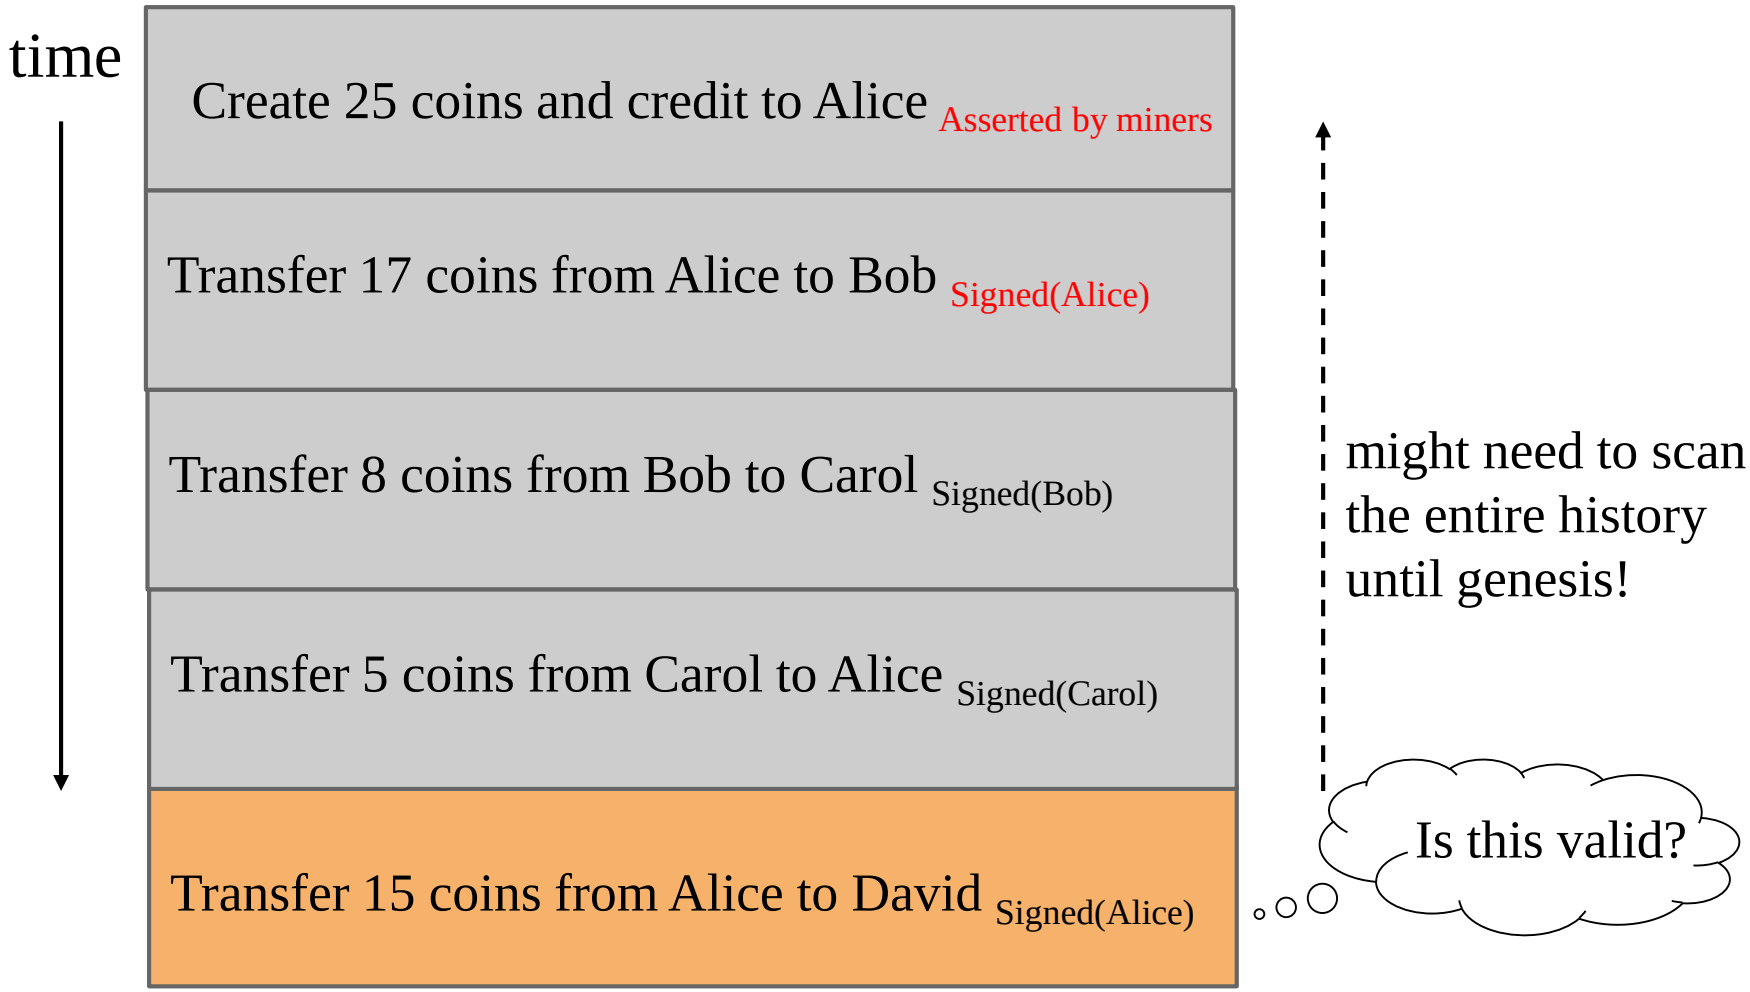

# Trans-based ledger with in/out pointer (Bitcoin)

- Each trans has inputs /outputs
  - inputs specifies source of coins; outputs the recipients of coins
- Easy to check if a transaction is valid (owner has sufficient coins?)

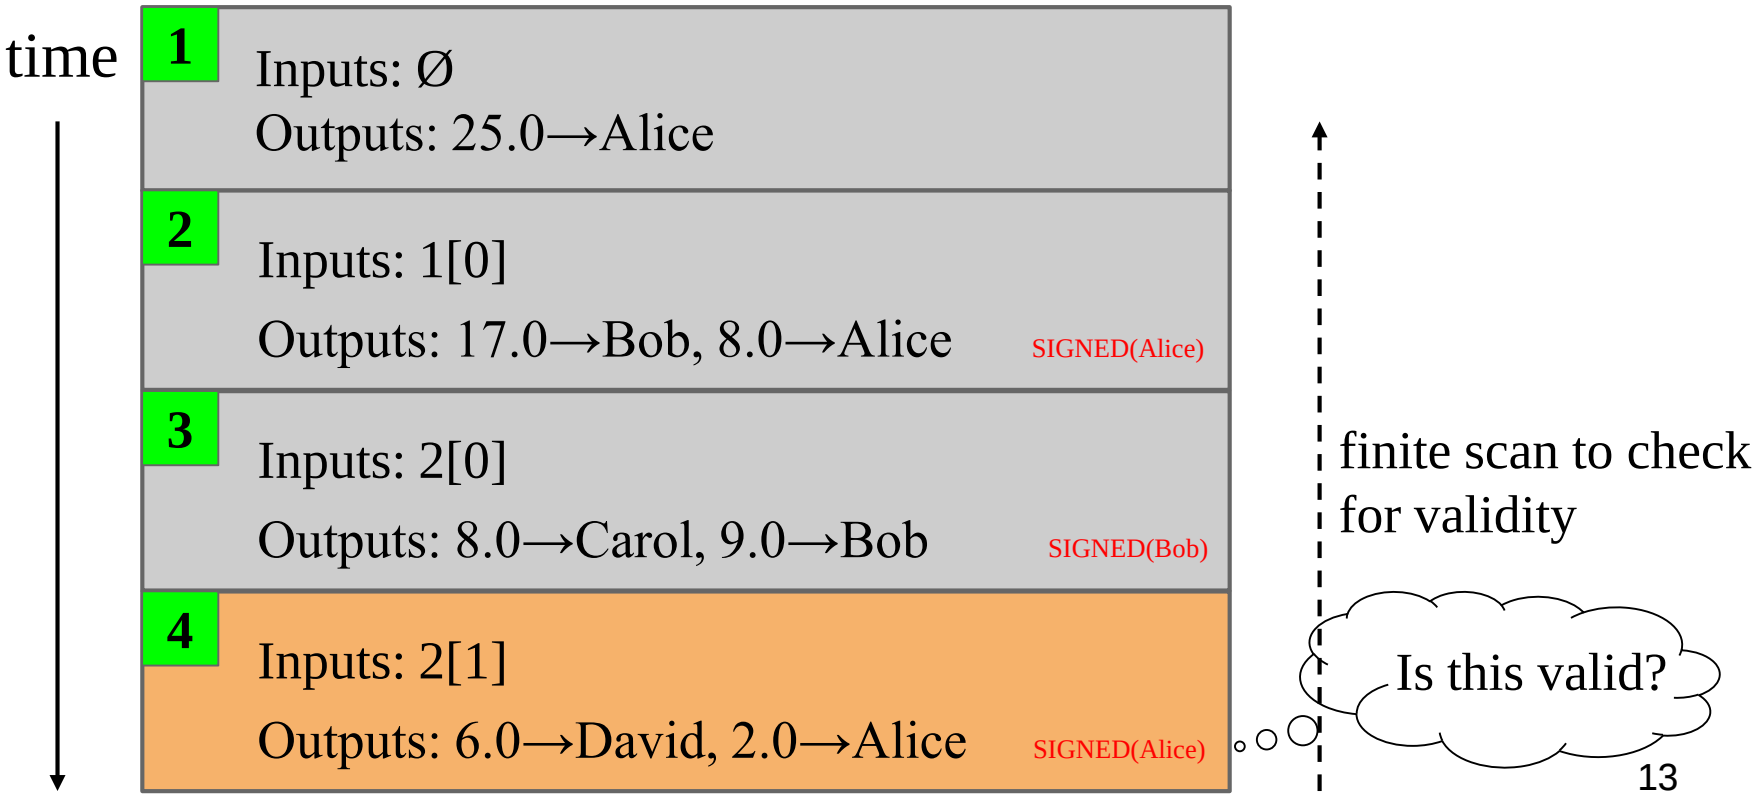

# Input/output link of transactions

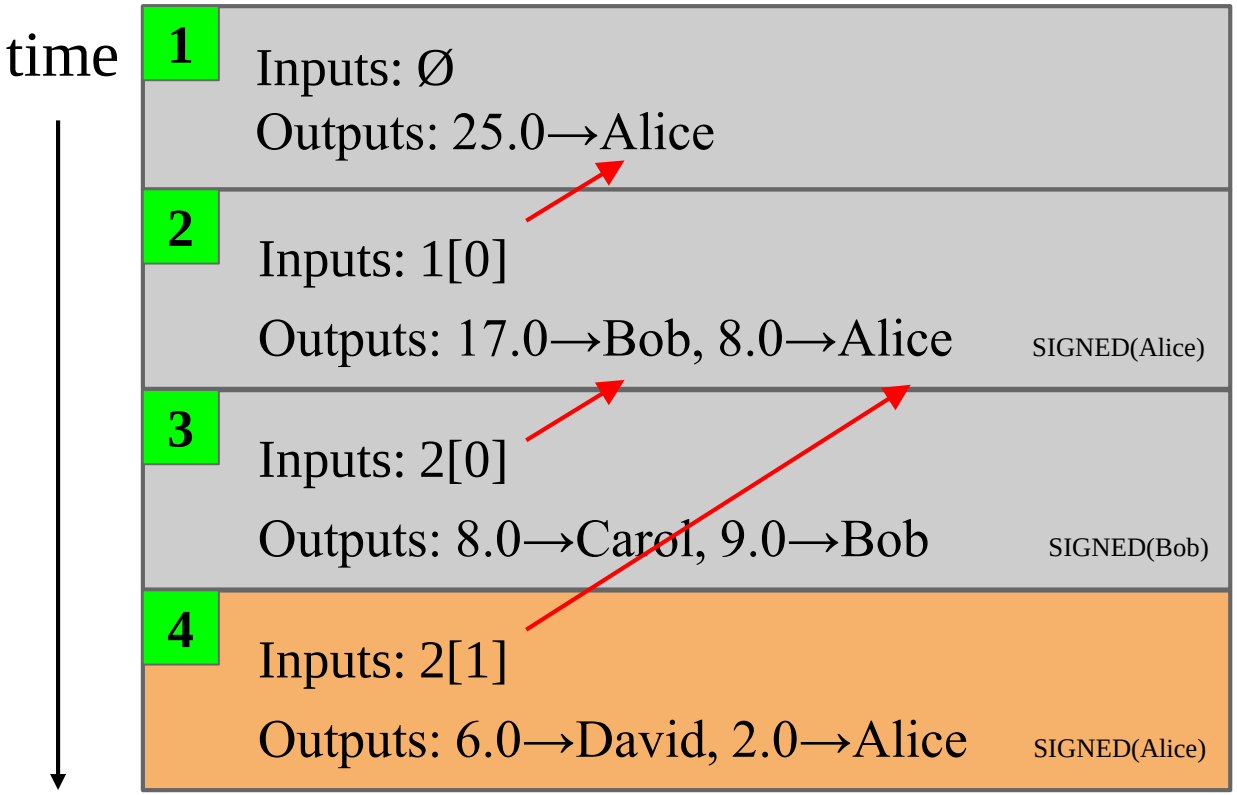


we implement this with hash pointers

SIMPLIFICATION: only one transaction per block

# A transaction with change:
## input value > transfer-value

- Alice has 9 coins and transfers 6 to Carol, and Alice still has 3 coins left
- The transaction has two outputs: one for transferring to Carol and the other for transferring back to Alice
- The total inputs always equal to the total outputs of a trans

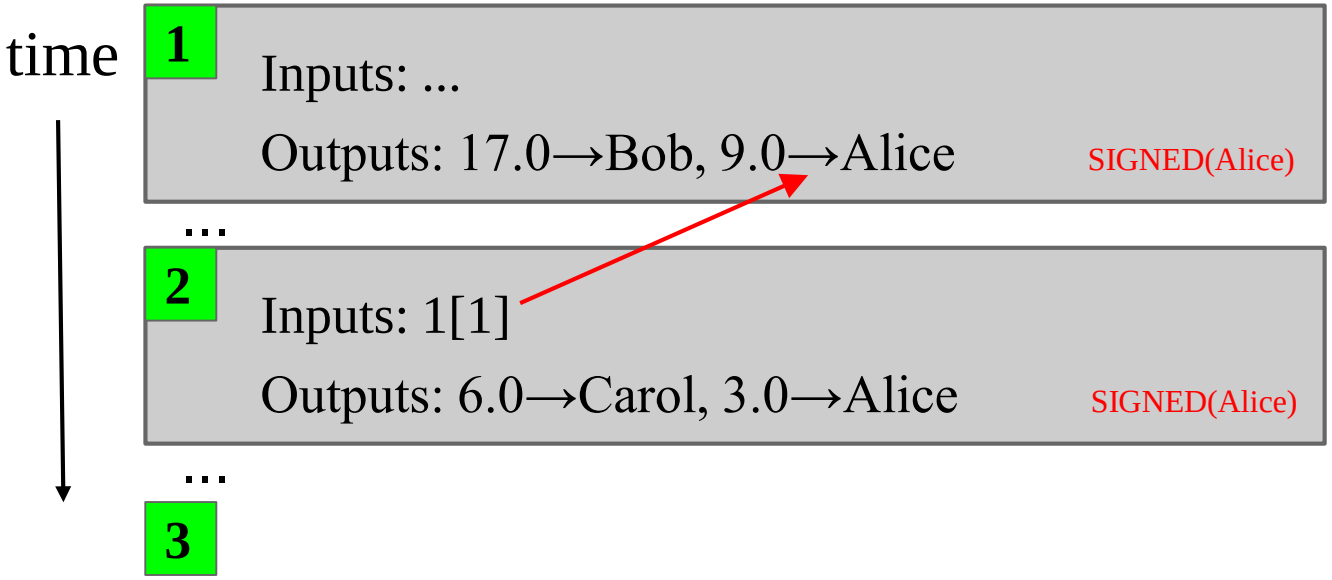

# Joint payment

- *inputs* can come from multiple sources
  - the transaction needs to be signed by all input owners

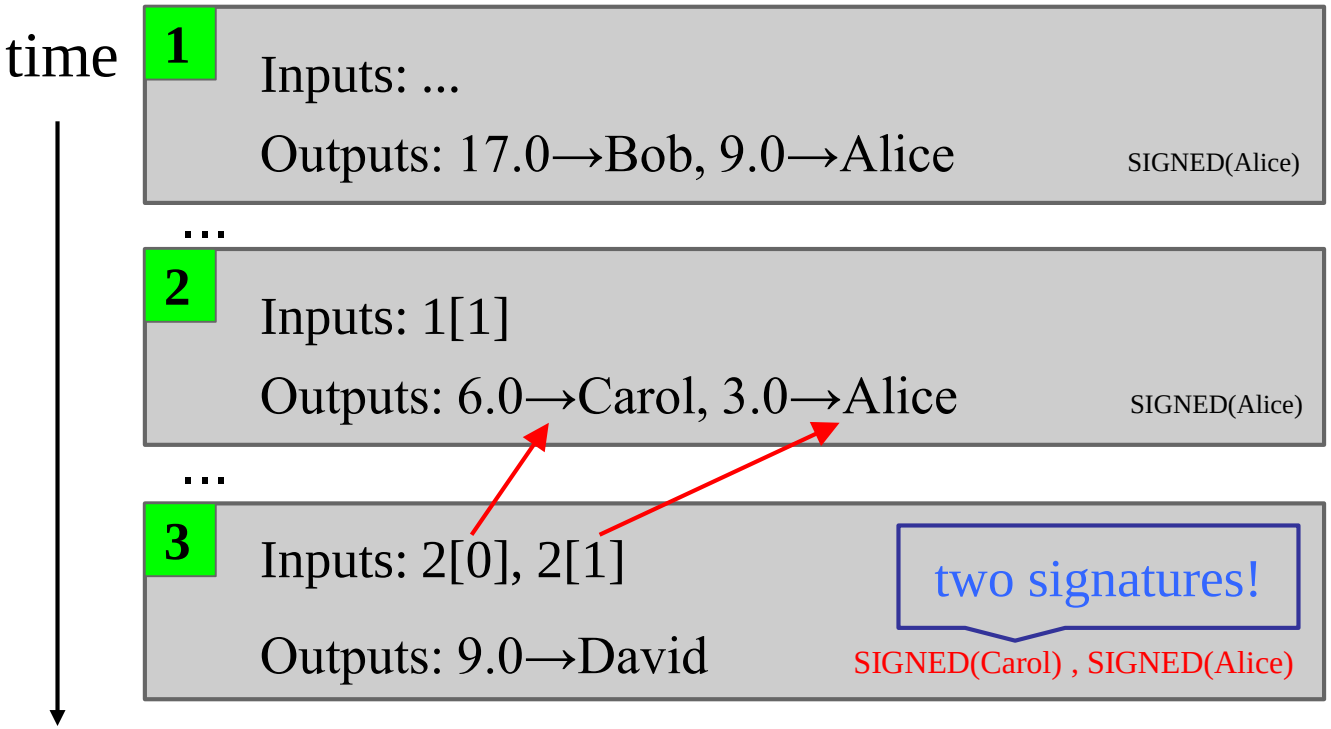

# Merge multiple outputs

- Merge outputs of multi-trans for the same owner
  - simplify the input of future trans, and
  - make it easy to verify balance of an owner
- The system can do auto-merge

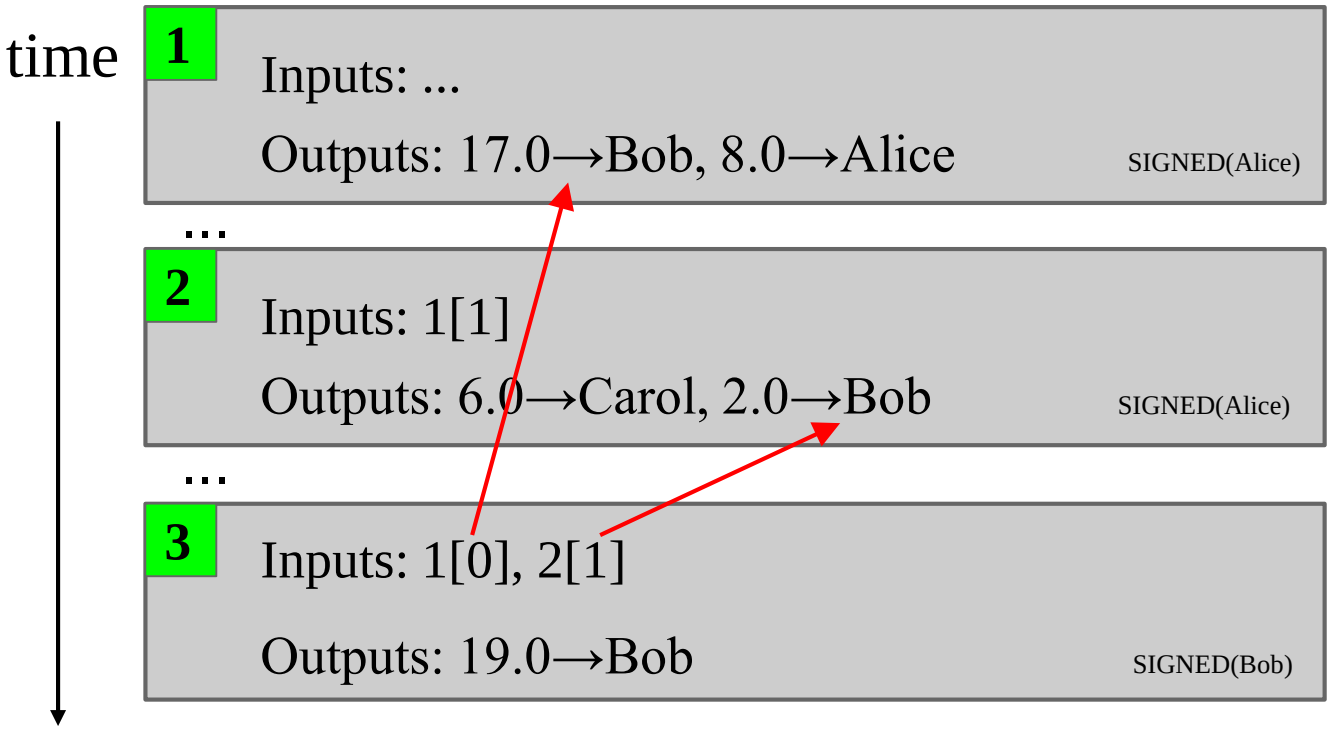

# Transaction syntax

metadata

input(s)

output(s)

```
{                    (transID)
    "hash":"5a42590fbe0a90ee8e...b8b6b",
    …
    "size":404,
    "in":[
    {"prev_out":{                (prev. transID)
    "hash":"3be4ac9728a0823ca…80260",
    "n":0}
    "scriptSig":"30440..."}(signature - script)
    ],
    "out":[
    {
    "value":"10.12287097",    (output value)
    "scriptPubKey":"OP_DUP OP_HASH160
     69e02e18b5705a05dd6b28ed51776c
     OP_EQUALVERIFY OP_CHECKSIG"}
    ]              (public key of recipient - script)
}
```

# Coinbase transaction

metadata

input(s)

output(s)

```
{
    "hash":"5a42590fbe0a90ee8e...b8b6b",
    …
    "size":404,
    "in":[
    {"prev_out":{      (null transID)
    "hash":"00000000000…000000",
    "n": 4294967295}
    "coinbase":"..."}  (arbitrary)
    ],
    "out":[
    {                  (block reward + trans fees)
    "value":"12.52287097",
    "scriptPubKey":"OP_DUP OP_HASH160
     69e02e18b5705a05dd6b28ed51776c
     OP_EQUALVERIFY OP_CHECKSIG"}
    ]
}
```

# Demo: block, transaction in blockchain

- Demo at https://www.blockchain.com/explorer

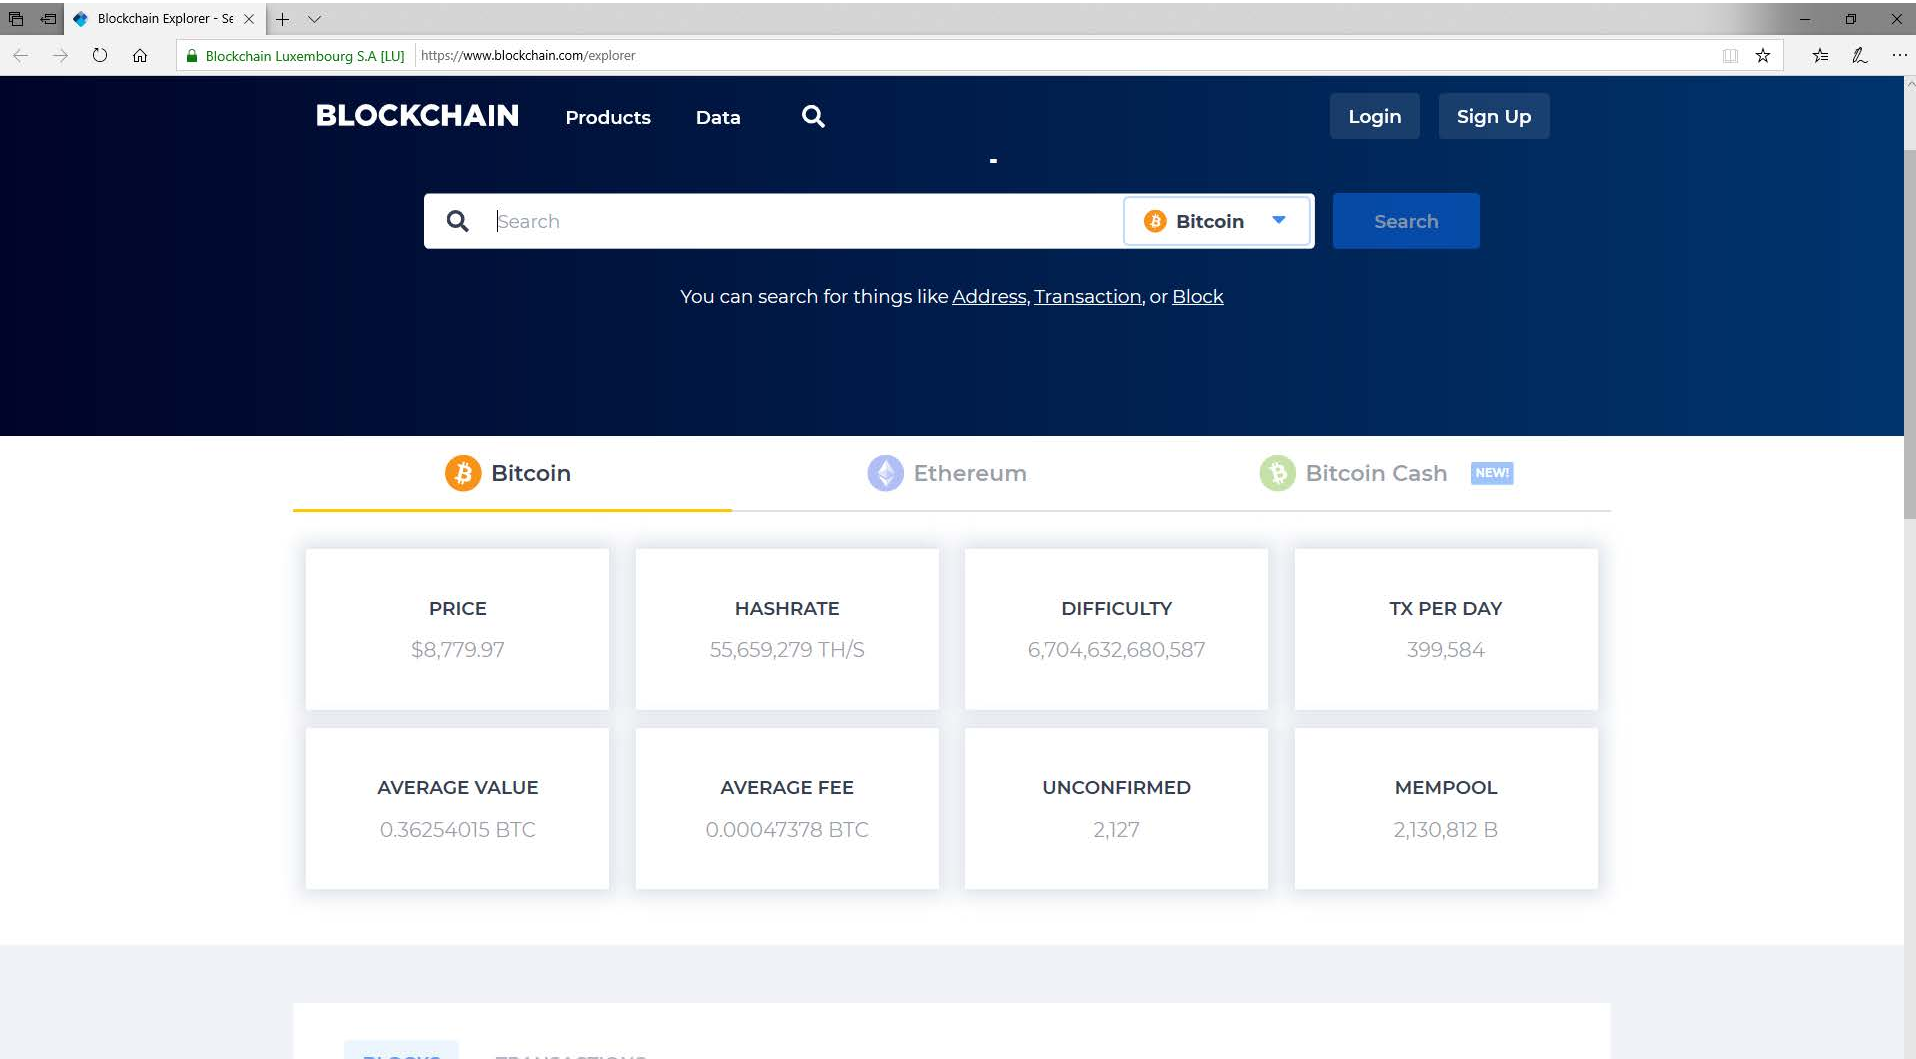

# Data structure of block: Chain of blocks

- Each block contains a set of verified transactions

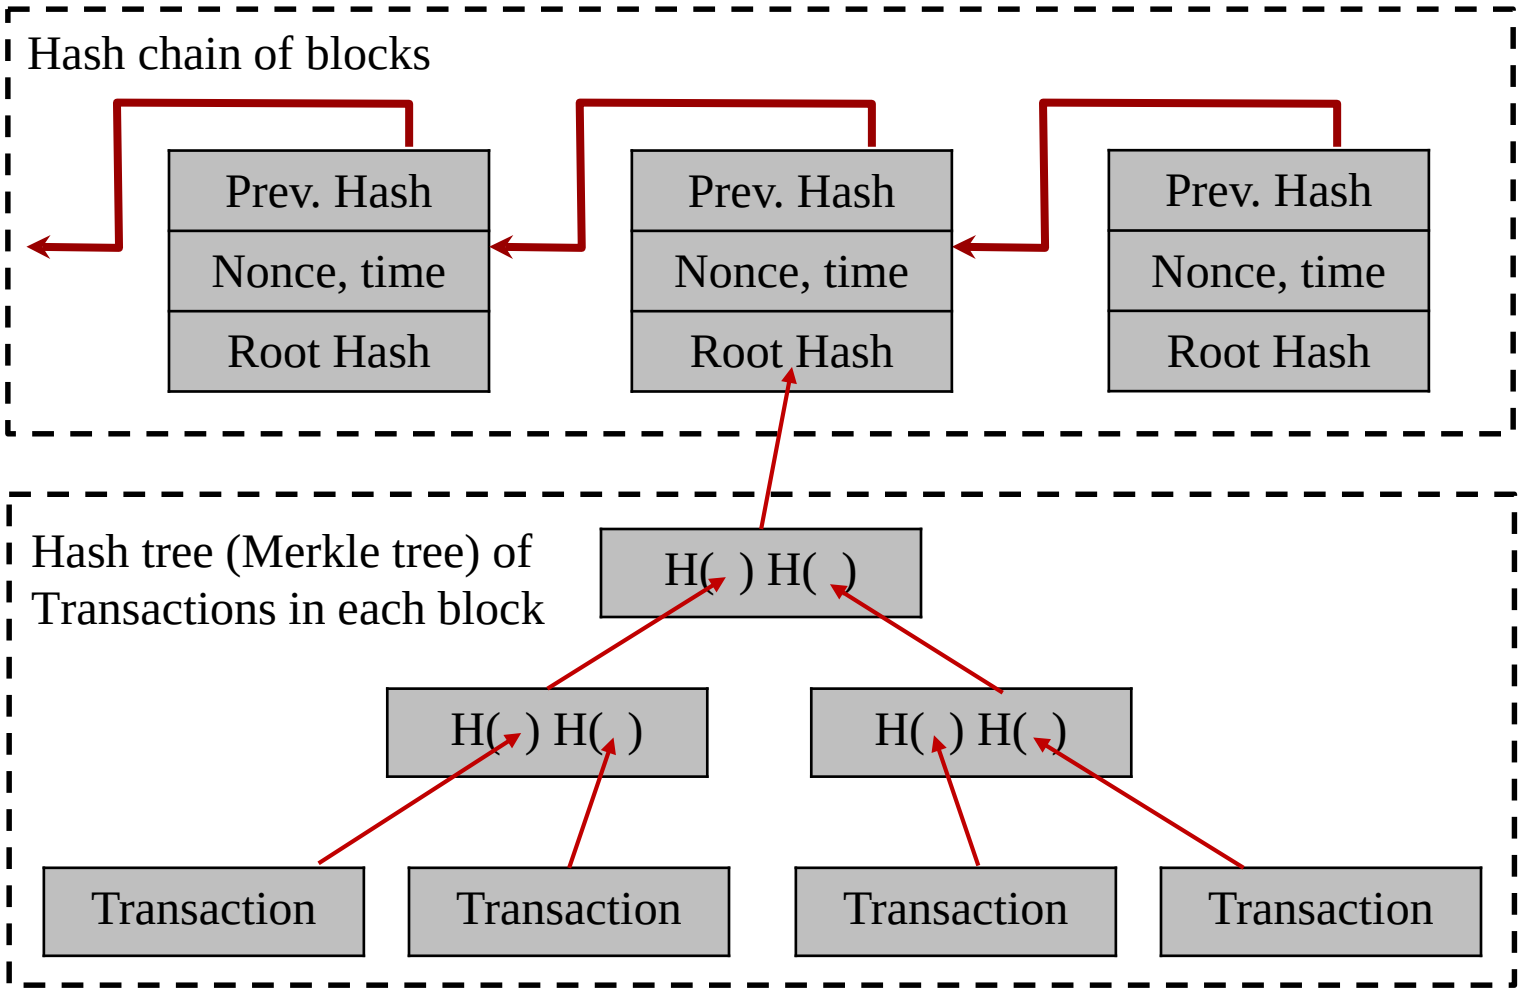

# Bitcoin block syntax

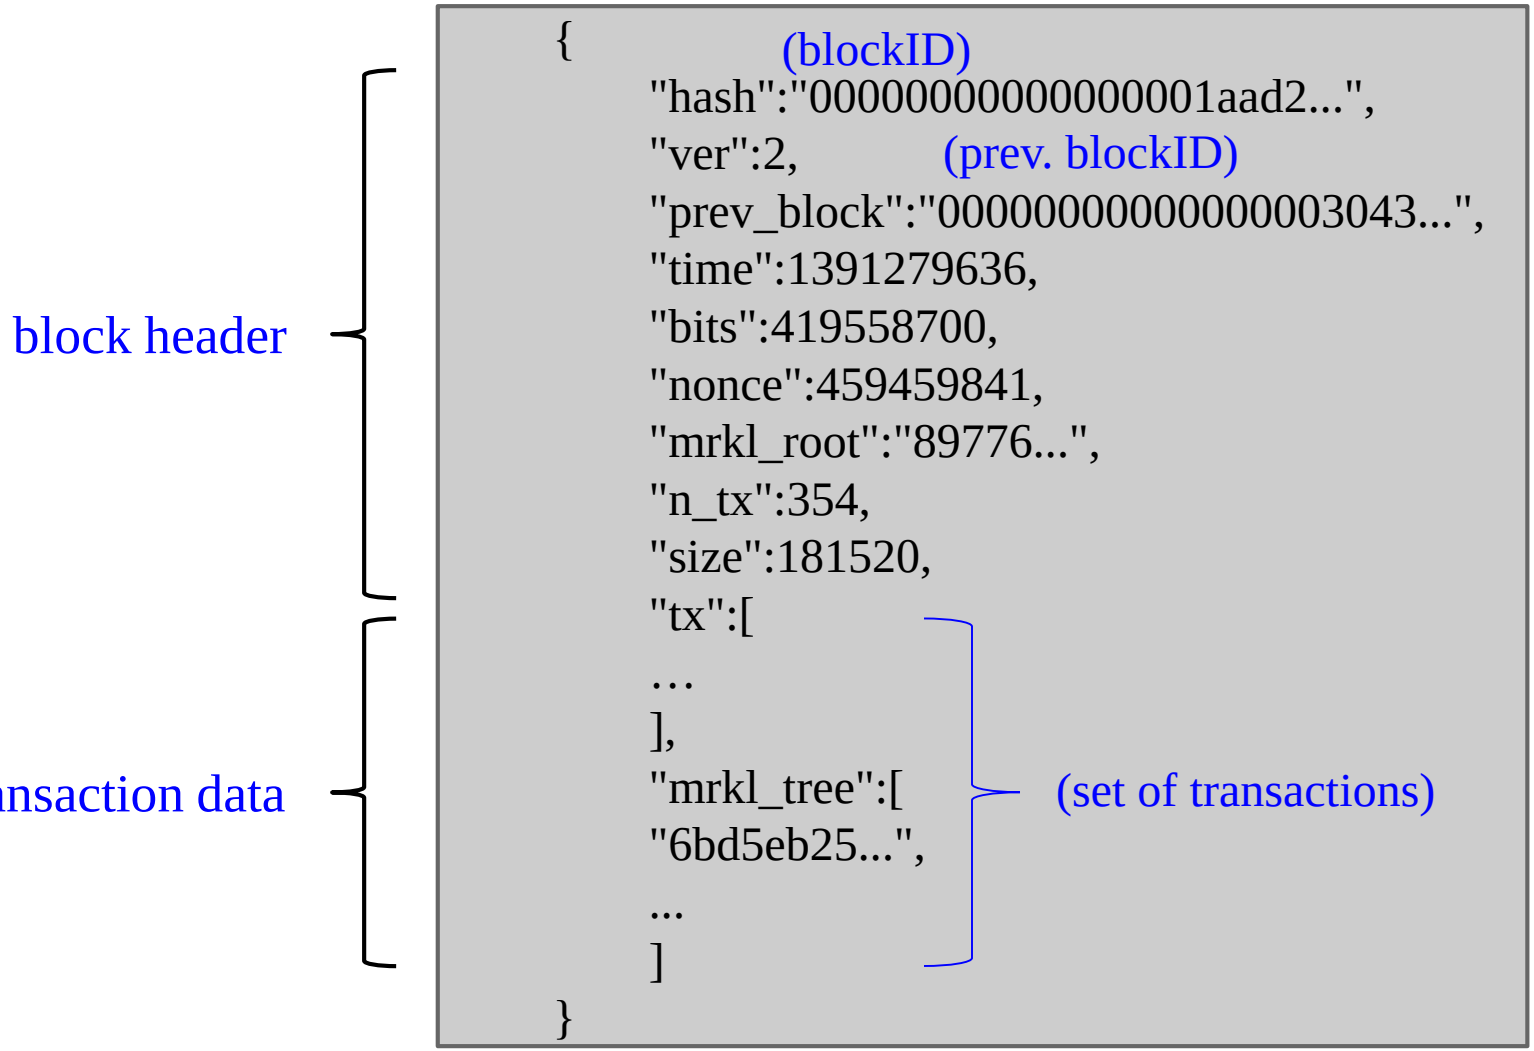

# Immutability and append only of blockchain

- Impossible to alter any transactions in the blockchain:
  - each node keeps a copy of the chain locally and all copies are consistent
  - each transaction is signed and verified

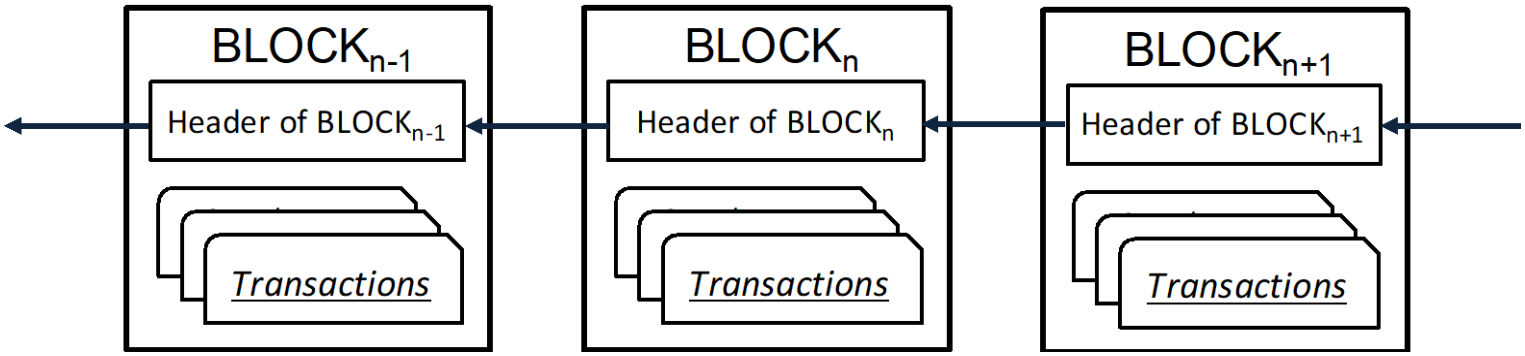

# Inconsistency of trans pools at nodes

- Each transaction is broadcast to all nodes and nodes have different sets of trans due to network delay
- Each node selects a subset of trans from its local pool, verifies them and competes with other nodes to solve PoW
  - A node broadcasts a new block if it successfully solves PoW before others
- A new block is accepted by a node if it builds the next block upon this block

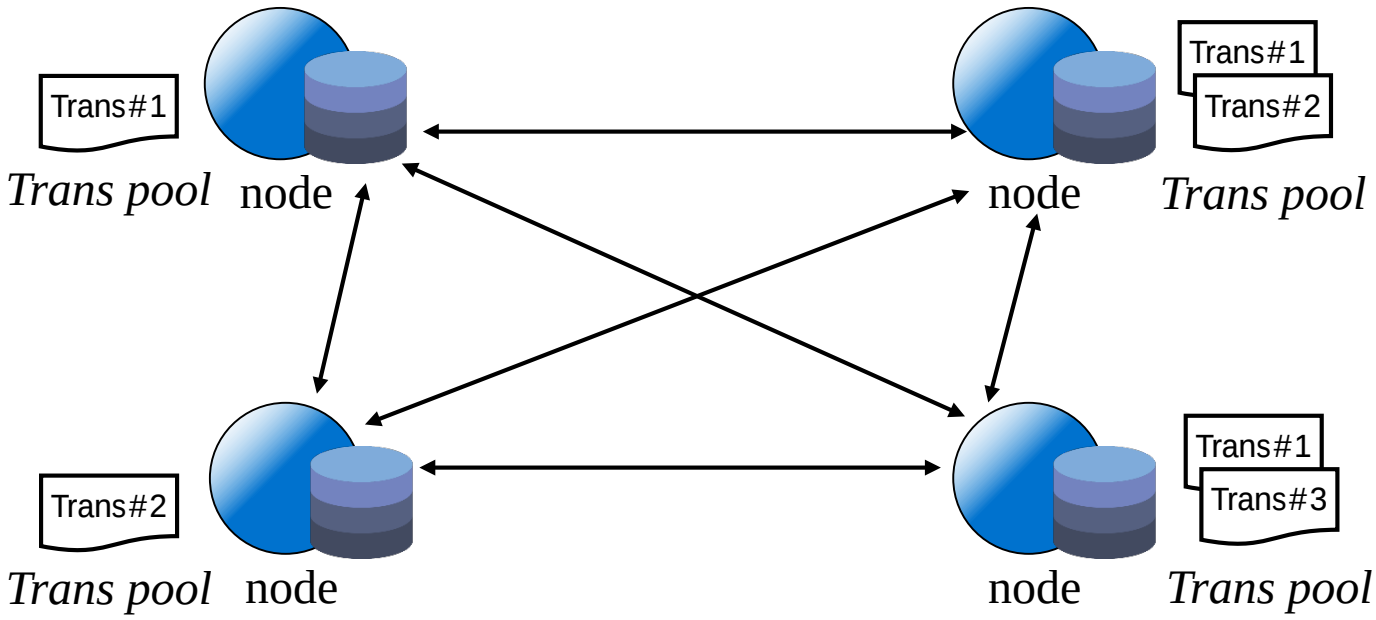

# Distributed consensus: Block mining

- Each miner (i.e., node) has a set of outstanding transactions it has received
- All miners execute a computationally–intensive process to decide which block to be extended

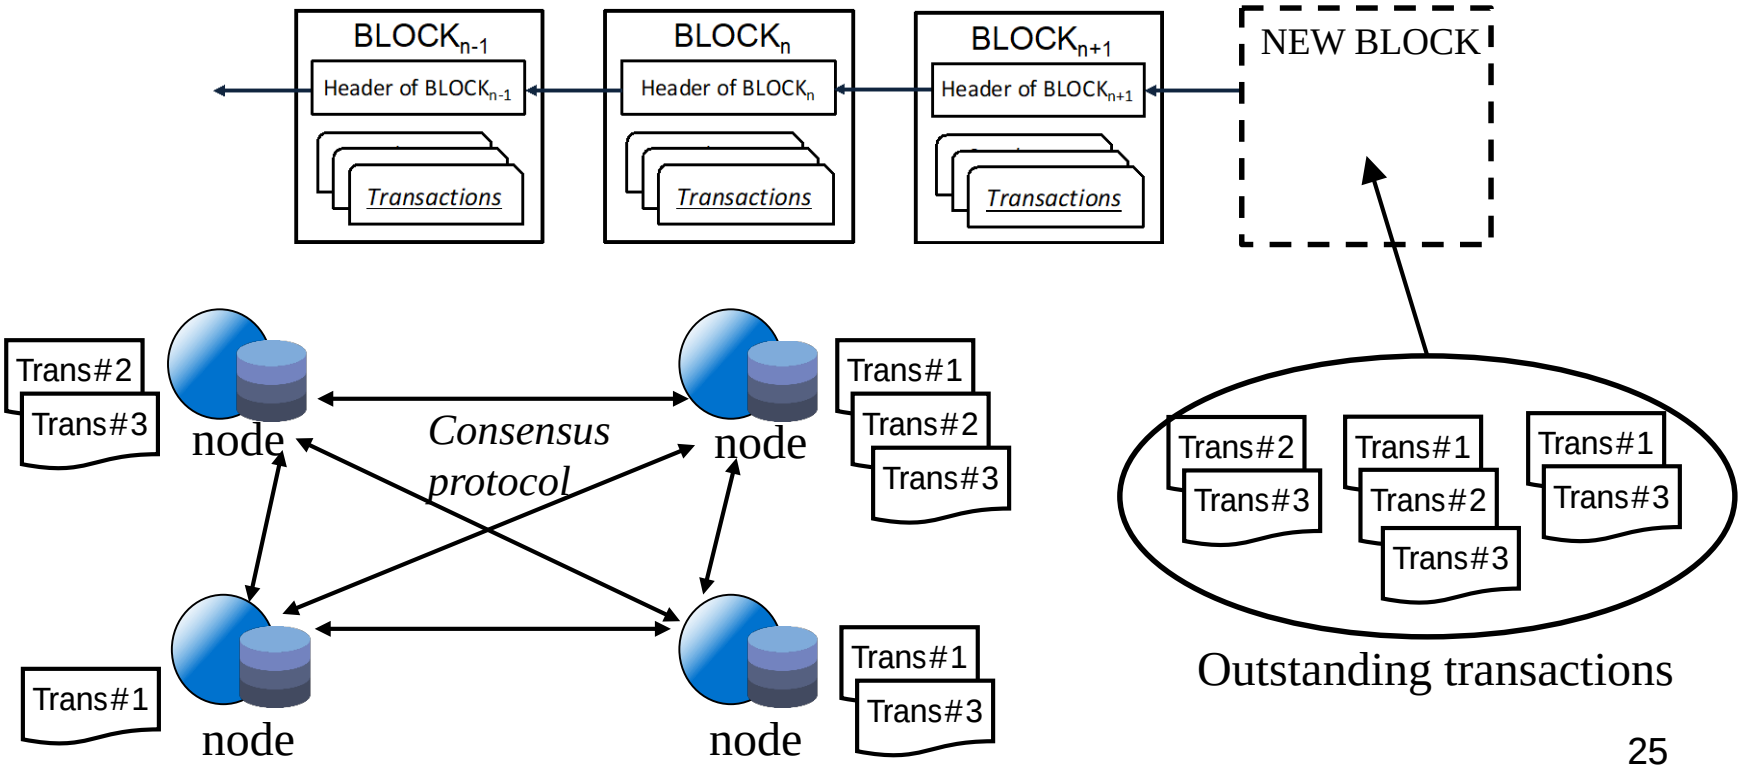

# Block mining: Proof-of-Work (PoW)

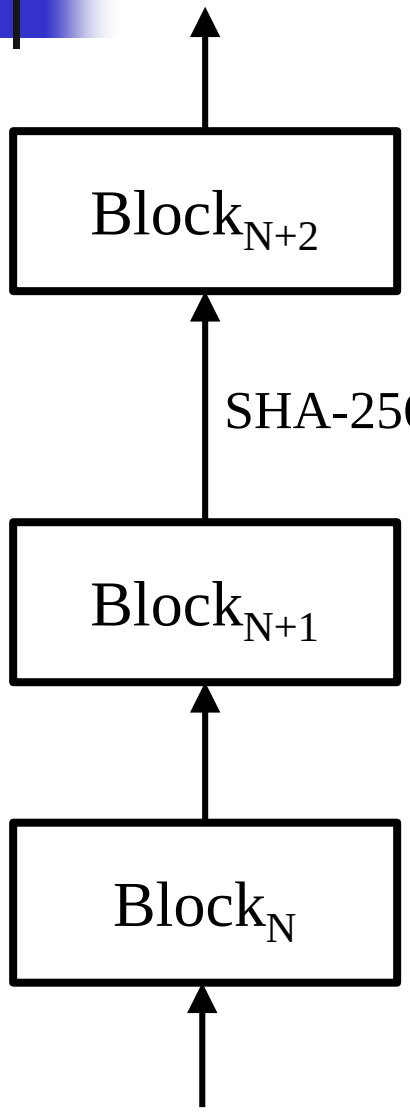


Mining difficulty

SHA-256^2(Block$_{N+1}$, X$_{N+1}$, ticket$_{N+1}$) = 0x0000000000000c67aa..

Mining problem: Find a ticket that produces a hash value less than target Z

SHA-256^2(Block$_{Prev}$, X$_{N}$, ticket$_{N}$) $\stackrel{?}{<}$ Z

software version, hash (Merkle-tree root) of trans, and current time

# Mining a new block: verify transactions and PoW

- Each miner picks a set of trans from its local pool & verifies them
- Computes the PoW and if successful:
  - link the block to the local chain, and
  - broadcast the block to the network

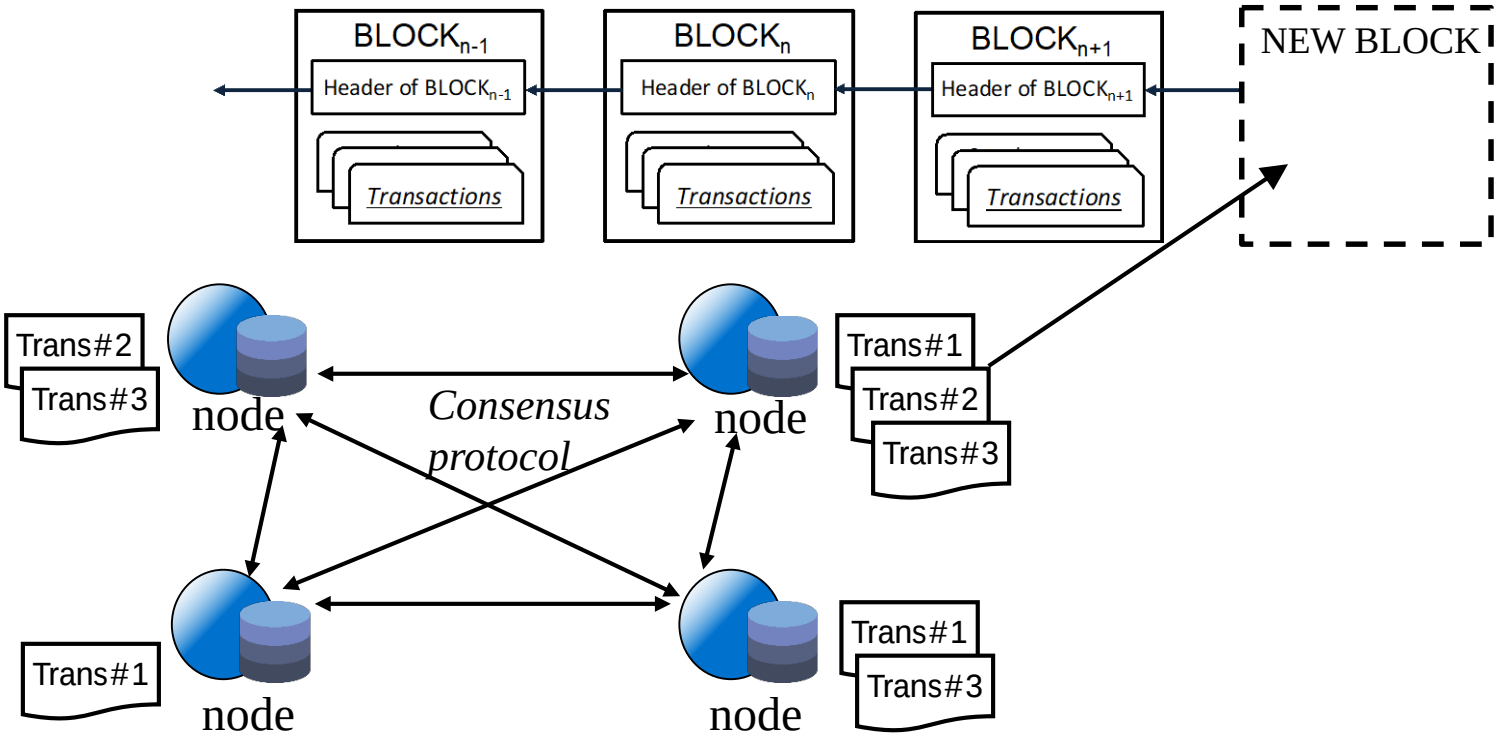

# Example: miners generate a new block

Each attempt has $16^{-3}$ chance of success

```
                                        Z = 0x000***...
Hash(Block 3 |  …  | 0xb9824 ) = 0x000c3f...
```

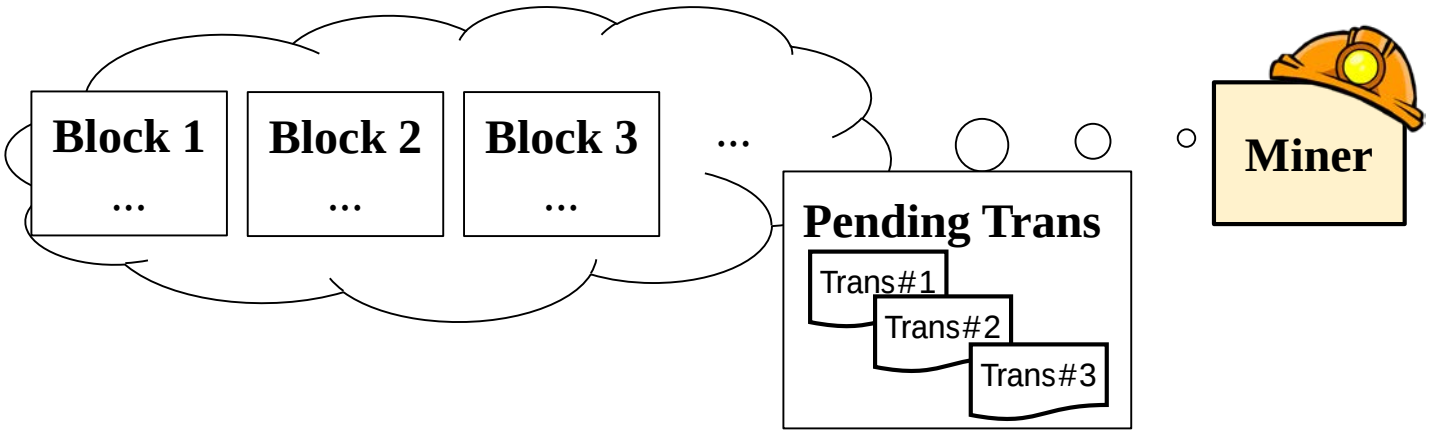

# What if a miner loses the competition?

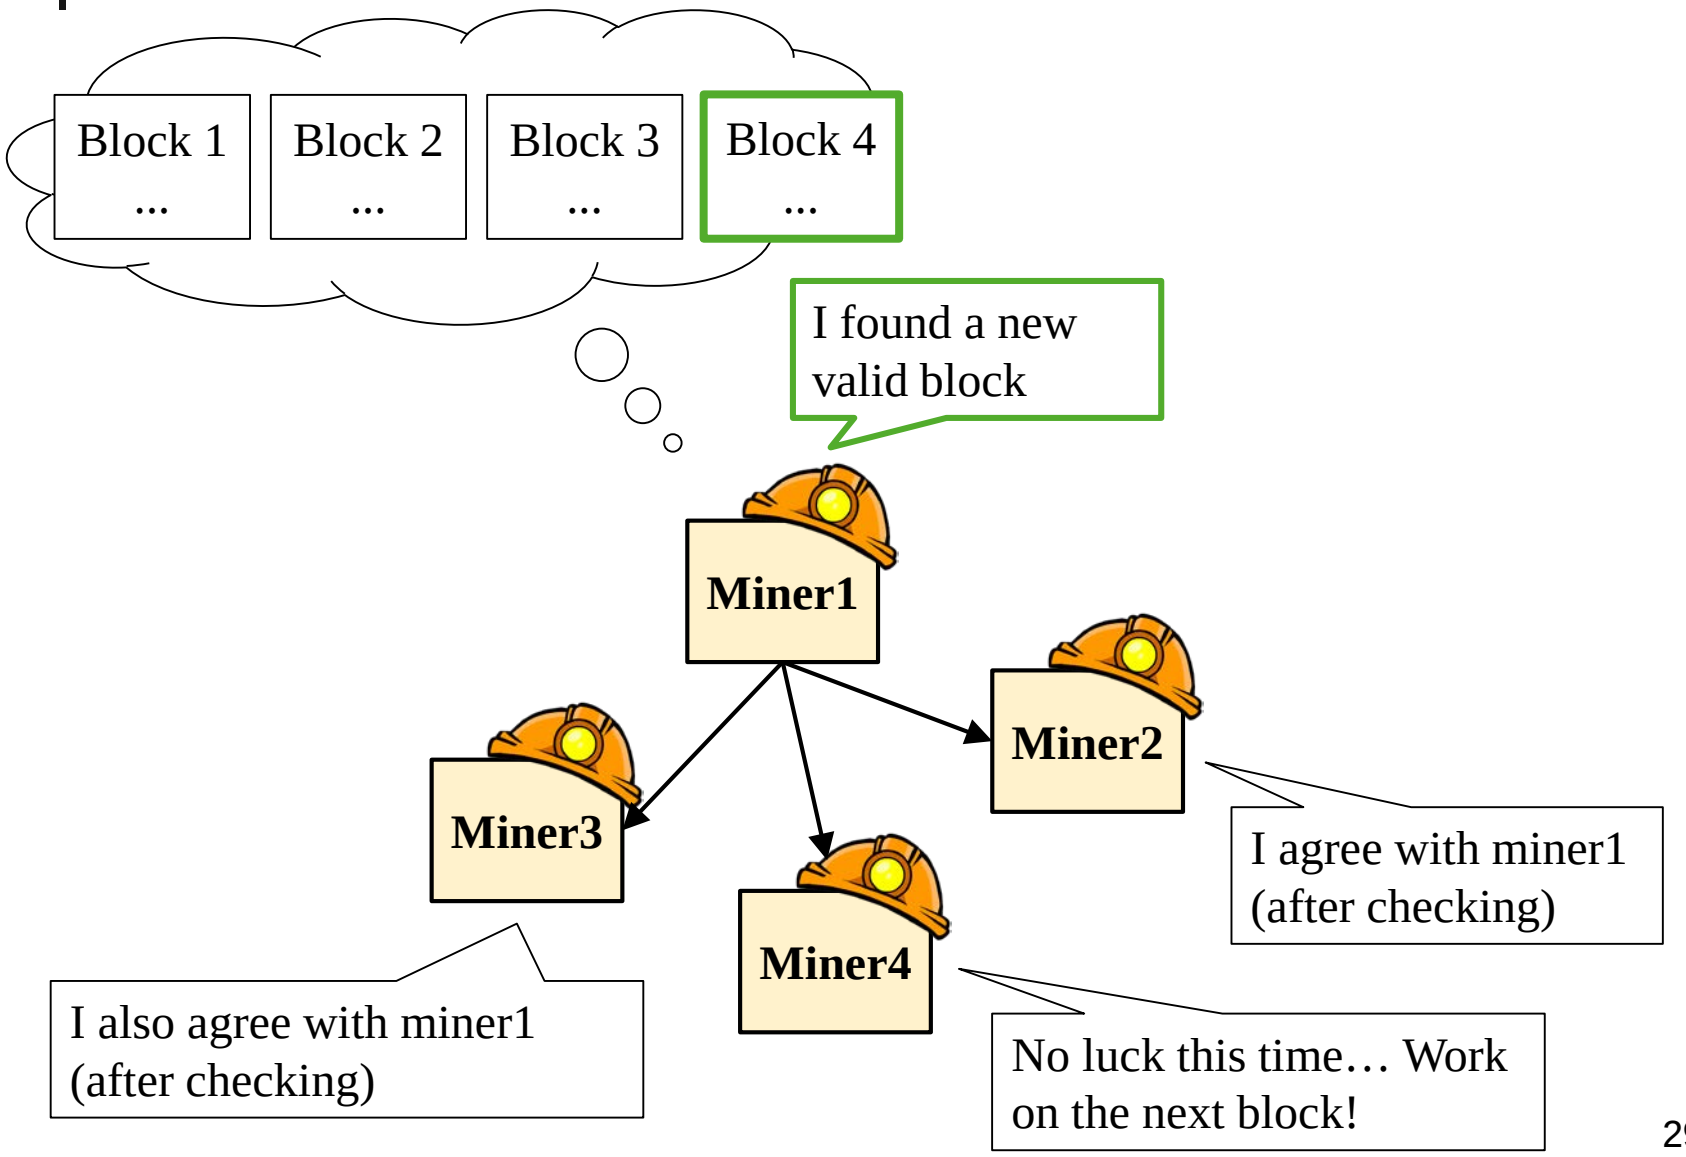

# What if two miners succeed simultaneously?

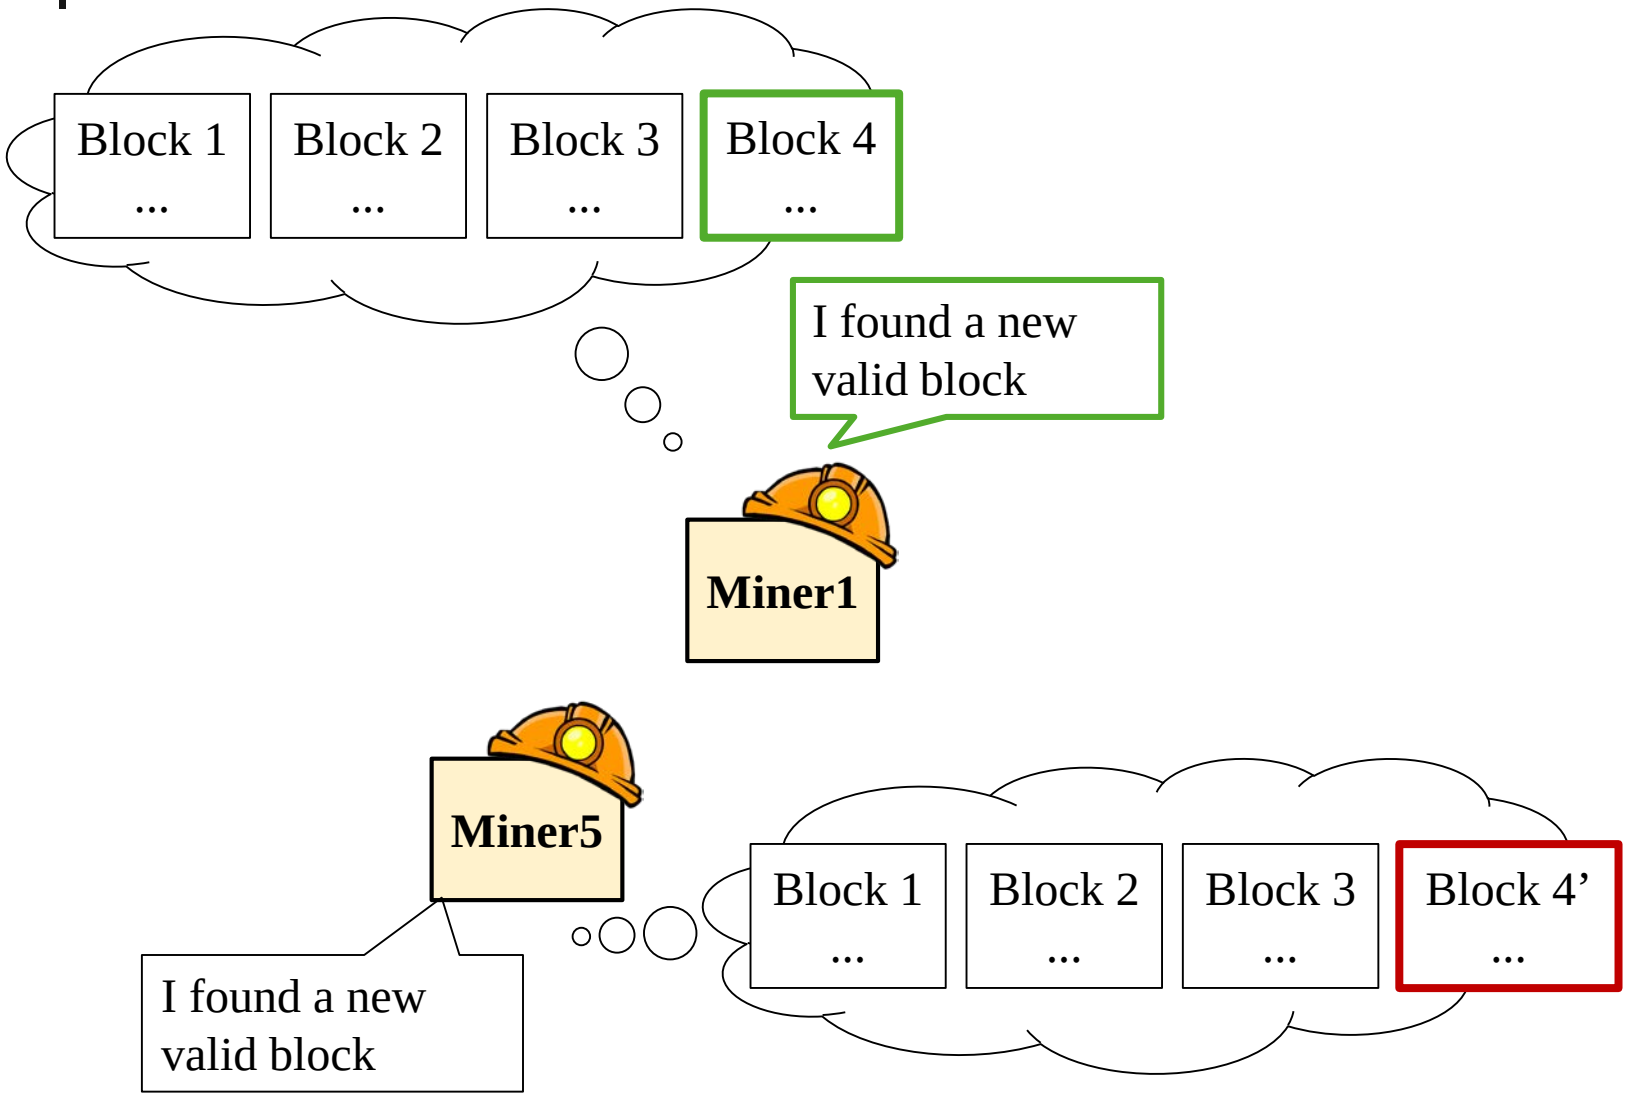

# Distributed consensus: Longest chain rule

- Two or more nodes may find a correct block simultaneously
  - a node that receives two or more new independent blocks will keep both blocks
    - The chain may temporarily have forks
  - it always works on (follow) a longer chain if there are multiple forks
    - Ties break arbitrarily
  - ~6 blocks ahead to confirm a transaction

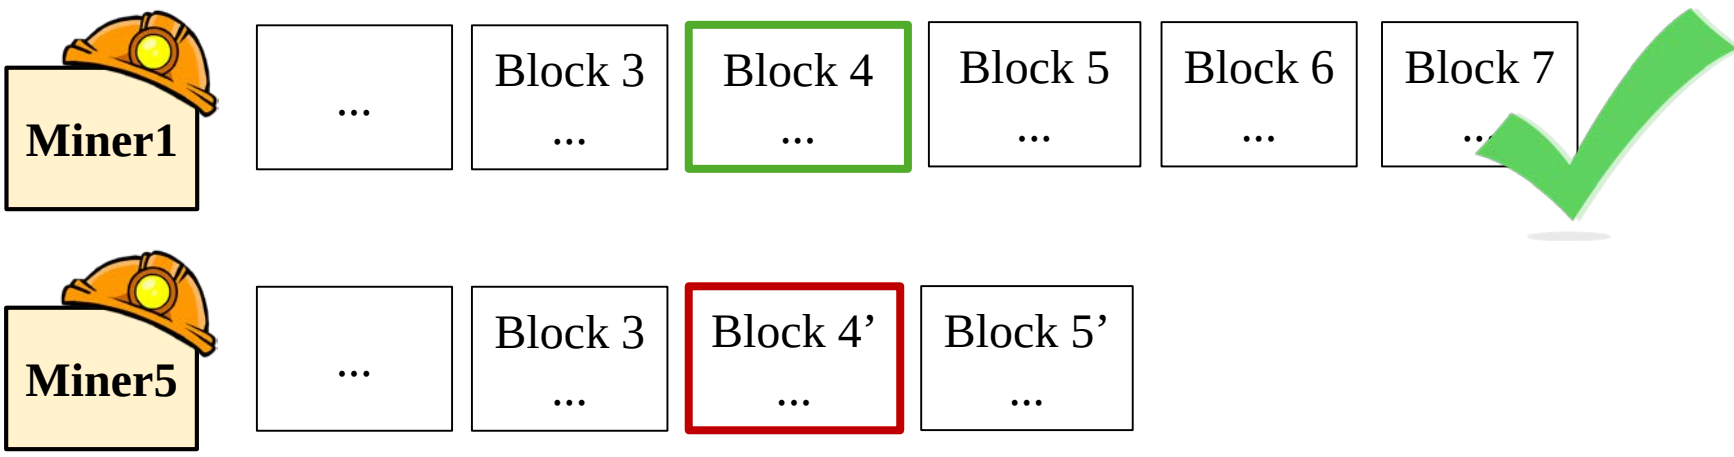

# Convergence to the same chain

- With the longest chain rule, all nodes eventually agree on the same blockchain
- Transactions of shorter blocks are put back to the pool
- How to reverse a trans that was already committed?
  - Do I see money credited to my account but later disappeared?

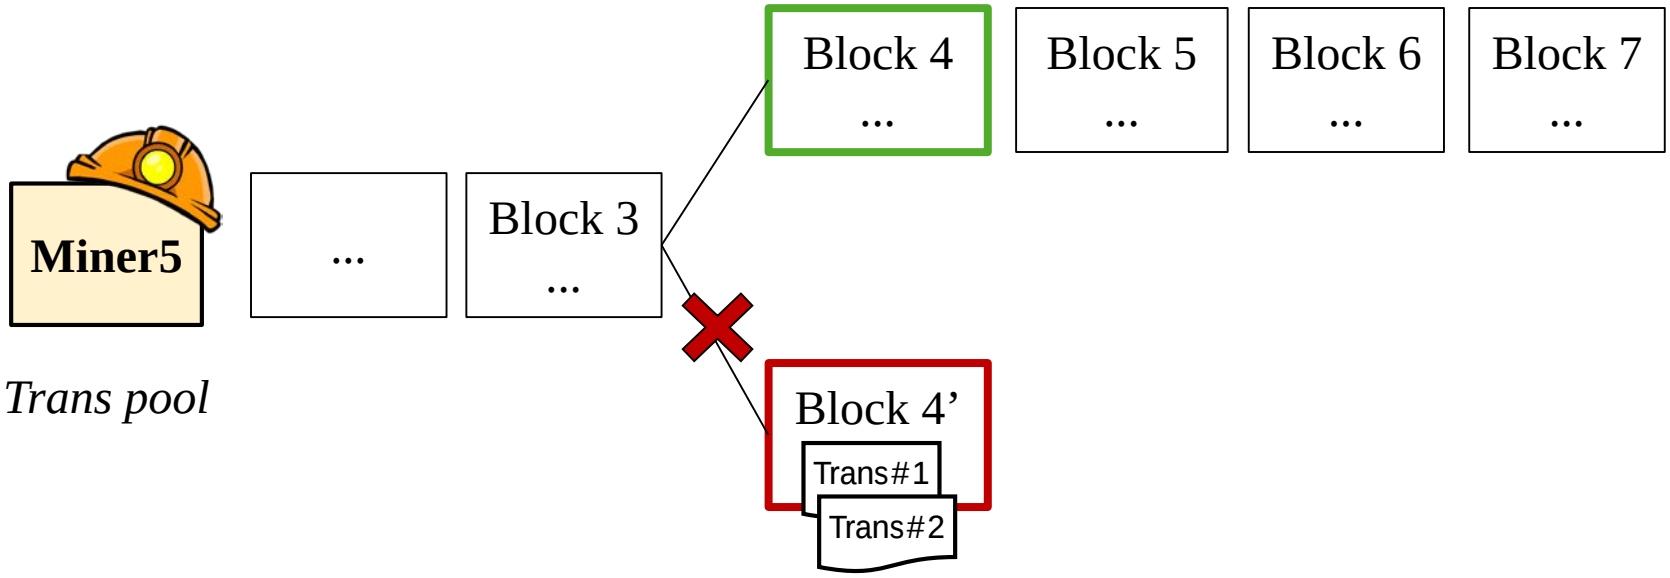

# When can a trans be confirmed in blockchain?

- There is no balance even written in the blockchain
  - The ledger is recorded as the history of transactions
  - When the trans of a lost block falls back to the pool, those trans are no longer in the chain
    - Fall-back trans take no effect, as if nothing happened
- When there are 6 (or more) new blocks grown after this trans, the trans can be regarded as “confirmed”
  - The funds in this transaction are then “committed”

# Impossible to fake a trans in blockchain

- Suppose a node made a fraud transaction and included in a block successfully
- This node has to continuously and successfully mine the next several blocks to make his faked block in the longest chain (even others can check out the fraud)
- But, the probability is very low:
  - suppose the node has *p*% of the total computing power…

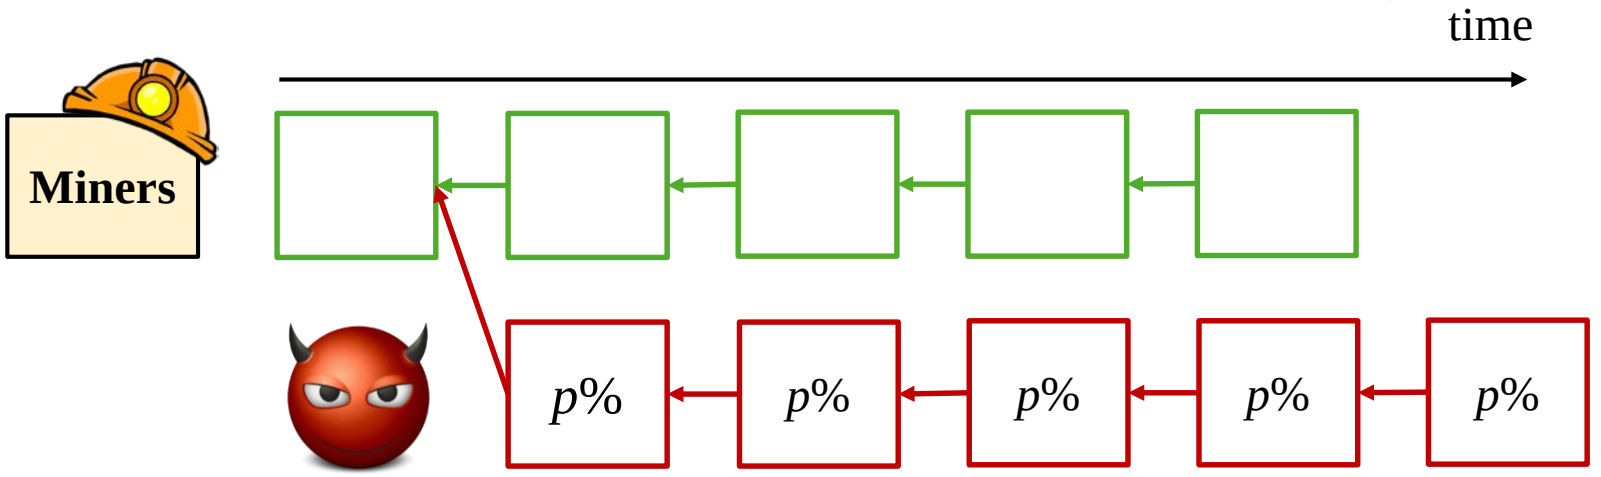

# What happens if a miner finds a faked trans?

- It simply doesn't follow the block for growing a new block
  - no reporting mechanism
  - note: no law-enforcement nor central-authority to catch the offenders in blockchain
- The owner of faked trans won't be able to keep up with the pace to generate subsequent new blocks
  - the block containing faked trans will be eventually discarded and the faked trans will never take effect in blockchain
- The counter-fraud in blockchain relies on the PoW and is based on the fact: nobody controls over 50% of the total computing power in the world

# Why PoW is essential?

- Spread out the time of nodes competing for generating new blocks in a wider range and with higher randomness
  - The probability for two miners to generate new blocks simultaneously is slim
  - Longer time for PoW makes network delays insignificant in winning out the competition among nodes
- Security reason
  - Prevent Sybil attacks
- Is it possibly to develop a decentralized consensus protocol without using PoW?
  - BFT (Byzantine Fault Tolerance) protocol
  - Proof-of-Stake

# Incentives for miners

- Block Rewards:
  - creator of a new block gets to include a special *coinbase transaction* in the block
    - The creator (typically itself) can choose a recipient address of this trans
- Transaction Fees:
  - a transaction's output value can be made less than the input value, leaving a transaction fee for the block creator
    - purely voluntary, like a tip
    - transaction fee becomes increasingly important, as block rewards start running out
- Where is Nakamoto's said 1M coins coming from?

# Maximum number of coins

- Coins are only generated through block mining
- The block reward is cut in half every four years
- Originally, 50 BTC/block; but today, 12.5 BTC/block

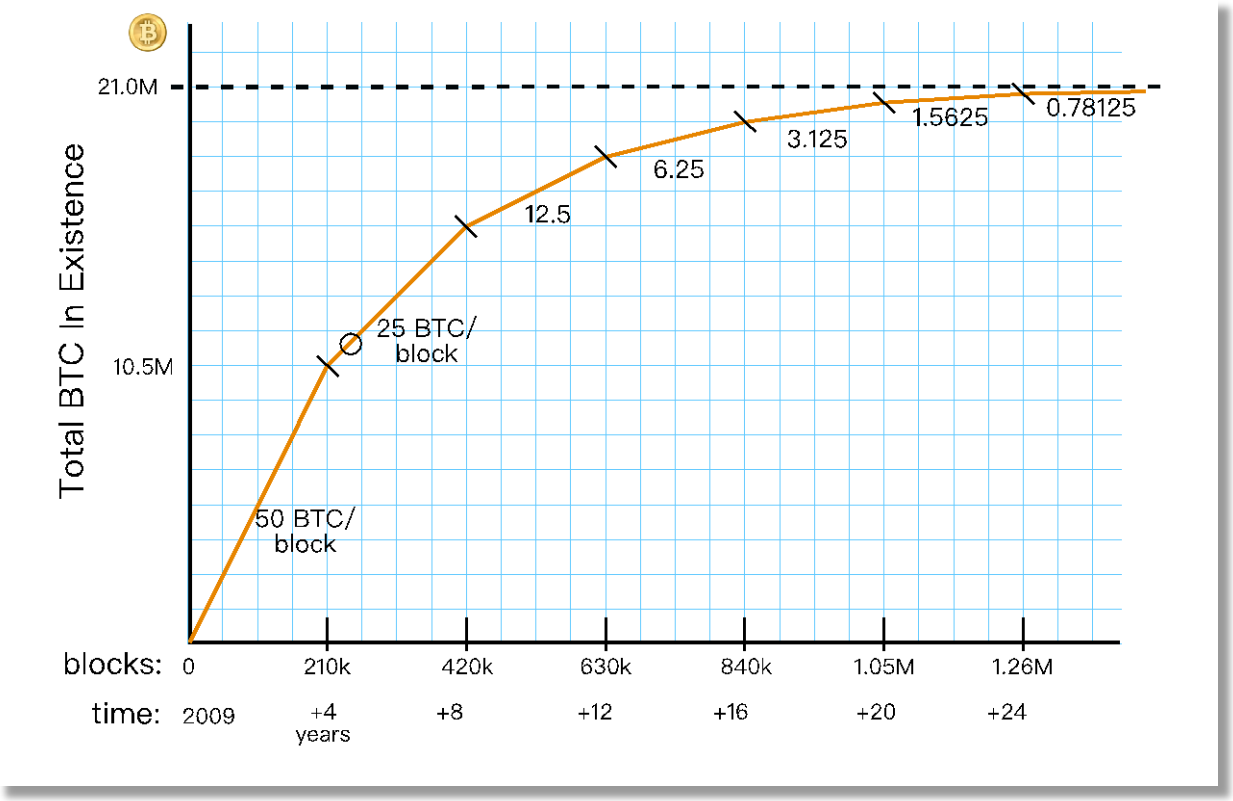

# Total number of coins is capped by 21M

- The number of blocks per 4 year cycle:

```
6 blocks per hour *
24 hours per day *
365 days per year *
4 years per cycle = 210,240 ~= 210,000
```

- Sum the block rewards for all years …

```
210,000*(50 + 25 + 12.5 + 6.25 + 3.125 + ... )
210,000*50(1 + 1/2 + 1/4 + 1/8 + …)
210,000*50*2 = 21 million
```

# Throughput of transactions

- Average time between blocks ≈ 10 minutes
  - nodes automatically re-calculate the difficulty of PoW every 2016 blocks (about every two weeks)
  - adjust difficulty to meet 10-minute goal
- Blocksize is limited to 1M bytes/block
  - at least 250 bytes/trans
  - ~3,500 – 4,000 trans/block
  - ~7 trans/s
- Compare to VISA (2,000-10,000 trans/s), and PayPal (50-100 trans/s)
- How to increase the throughput of Bitcoin?

# Bitcoin wallets

- You don’t need to mine or run a full node to use Bitcoin
- Wallet are applications that permit easy management of Bitcoins
- Bitcoin wallet stores, protects, and allows use of *private key* to make transactions

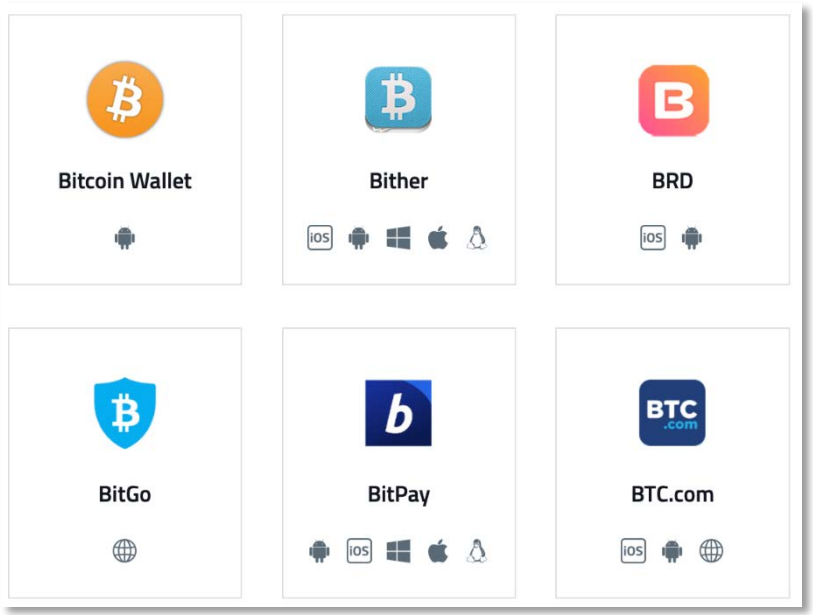


https://bitcoin.org/en/choose-your-wallet

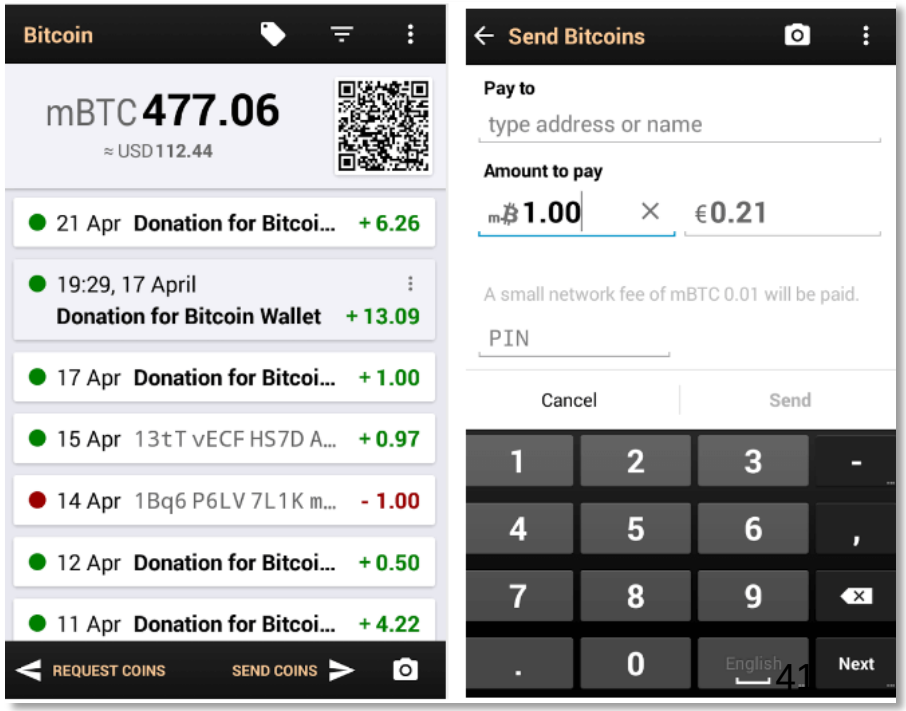

# Bitcoin payment

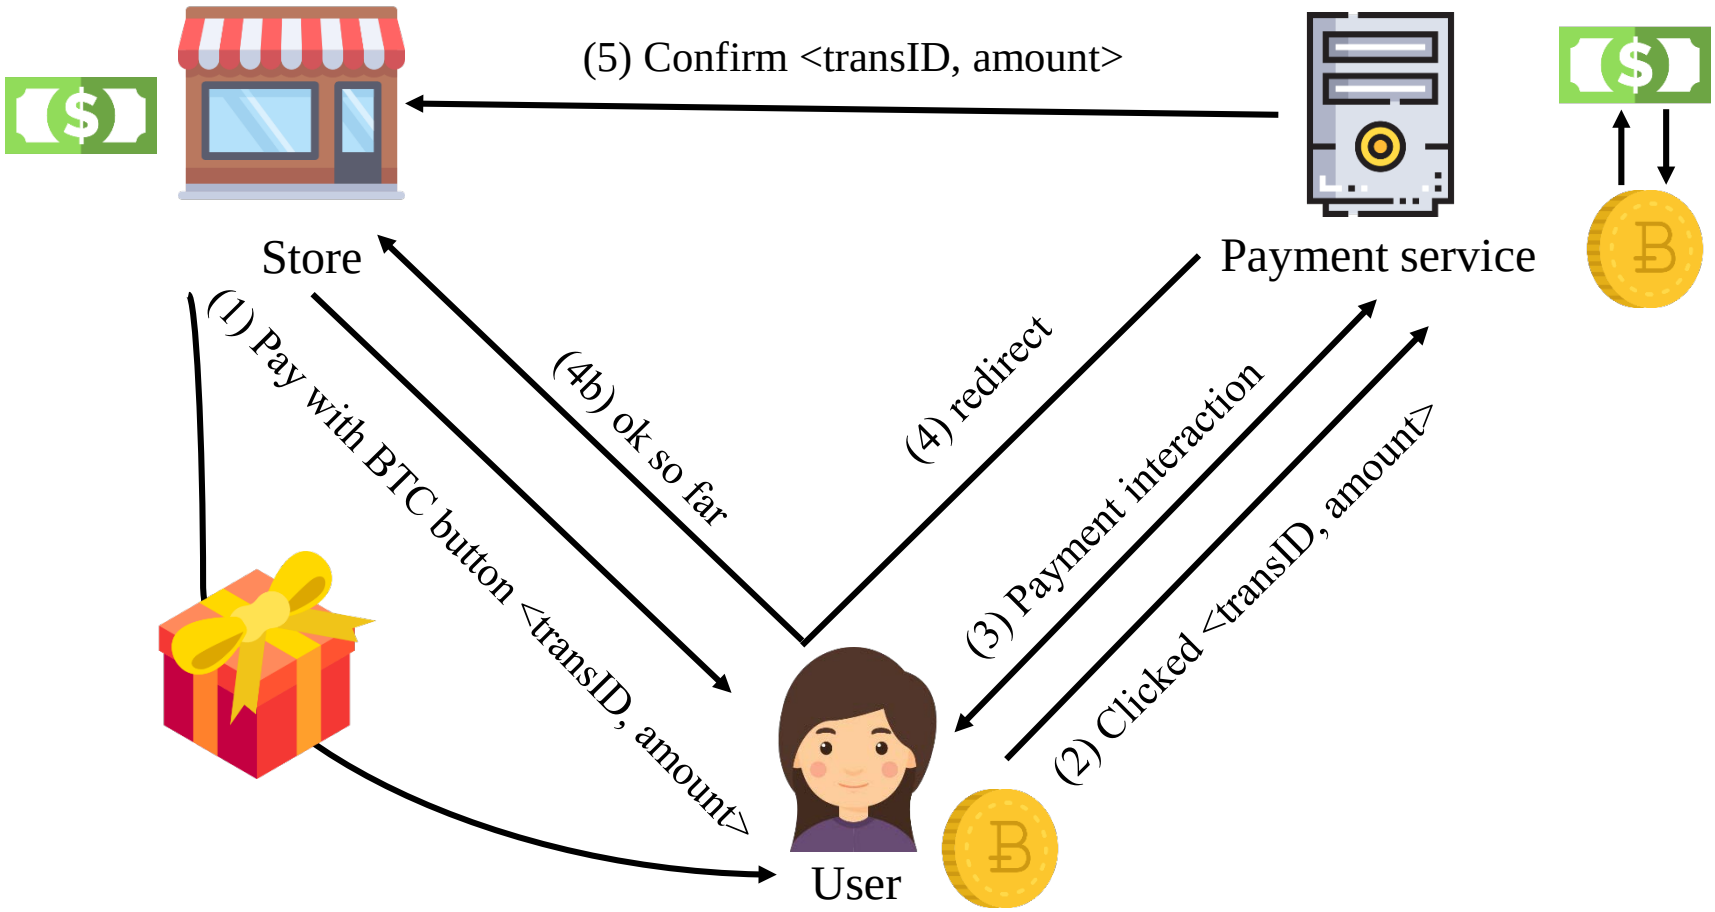

# Bitcoin exchange

- There are sites like *bitcoincharts.com* that show the exchange rate with various currencies
- Another option is to meet people to trade bitcoins in real life, such as *localbitcoins.com*

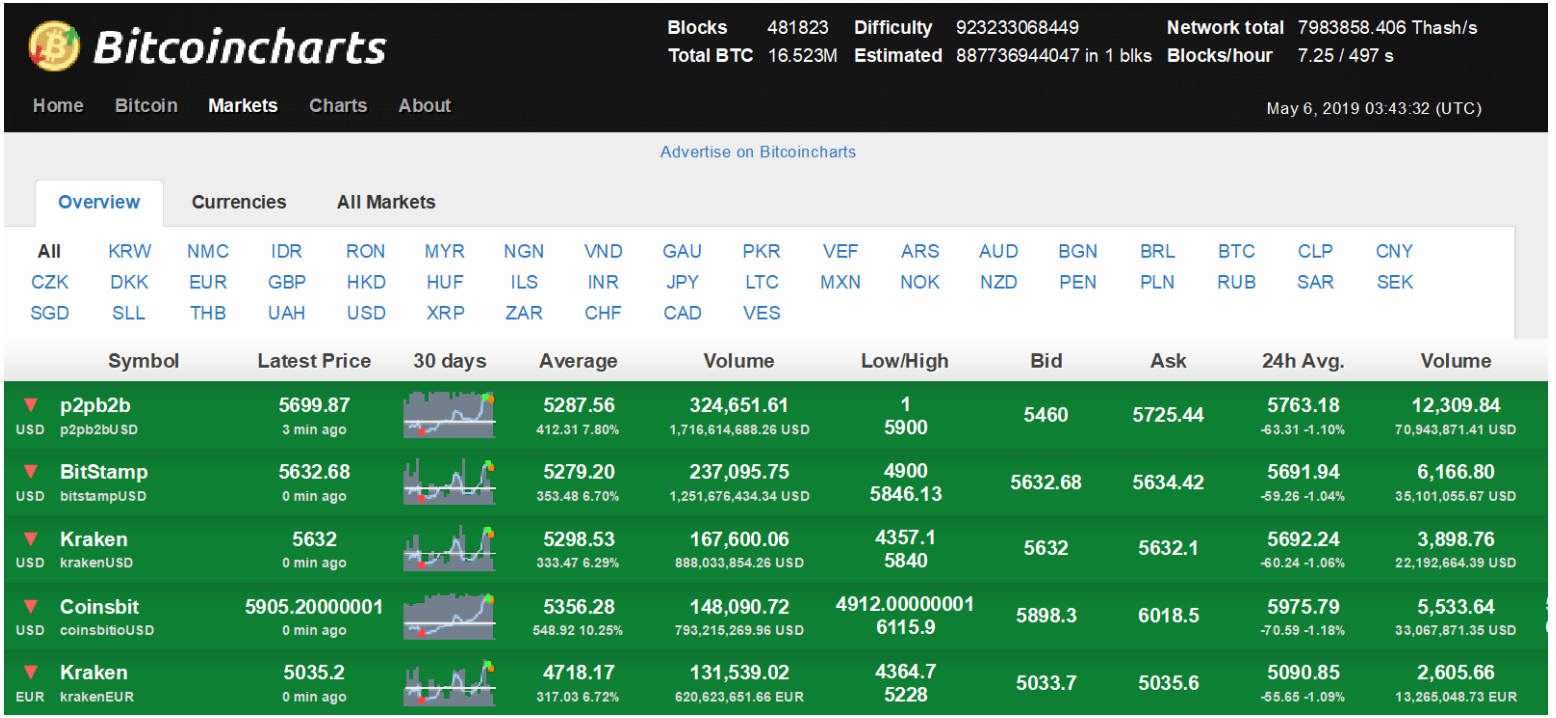

# Bitcoin's dark side

- Bitcoin has stimulated
  - Money laundering
  - Illegal marketplaces and dark web (e.g., Silk Road)
  - Ransomware
  - Theft of Bitcoin wallets
  - Rogue mining
    - E.g., ZeroAccess botnet

# Bitcoin's dark side


**Tor + Bitcoin = End-to-end anonymity for commercial transactions**

# Summary

- Bitcoin is a native application of blockchain technology
- The blockchain is maintained by a P2P network
  - each transaction is broadcast to the P2P network
  - miners verify transactions and generate new blocks to link to the chain
- The P2P network maintains the consistency of the blockchain via the longest chain rule
  - distributed consensus is enforced via PoW
- Blockchain technology can be applied to P2P environment where there is no central authority and no trust among the peers
  - Financial/banking sectors, insurance services, real-estate transactions, medical data sharing, etc

# References


- Joseph Bonneau, Andrew Miller, Jeremy Clark, Arvind Narayanan , Joshua A. Kroll, and Edward W. Felten. “SoK: Research Perspectives and Challenges for Bitcoin and Cryptocurrencies”, in Proc. of IEEE S&P 2015.
- Arvind Narayanan, Joseph Bonneau, Edward Felten, Andrew Miller, Steven Goldfeder. “Bitcoin and Cryptocurrency Technologies”, in Princeton University Press, 2016
- Satoshi Nakamoto. “Bitcoin: A PeertoPeer Electronic Cash System”
- Bitcoin Wiki, online at https://en.bitcoin.it/wiki/Main_Page
- Maurice Herlihy. “Blockchains from a Distributed Computing Perspective ”, 2018

# Exercises

1) Why is it impossible to make a fraud transaction in blockchain?
2) PoW costs a massive amount of resources. Why is it essential in blockchain? Can you replace the PoW by a protocol without heavy computational cost?
3) Why the max number of Bitcoins is capped by 21M?
4) The throughput current bitcoin system is around 7 trans/s, too small. Think about some ways to increase the throughput of bitcoin transactions, and discuss their pros and cons.
5) Think about an application that can use blockchain technology.

# Case Study: a P2P storage system using blockchain

- Explosive growth of digital data
  - fuelled up by e-health, e-commerce, smart cities, IoT, …
- Mismatch between supply and demand of data storage
  - a vast amount of under-used storages scattered all over the world
  - high demand from users looking for storage space
- P2P storage system:
  - utilize the unused storage space to form a huge global storage system

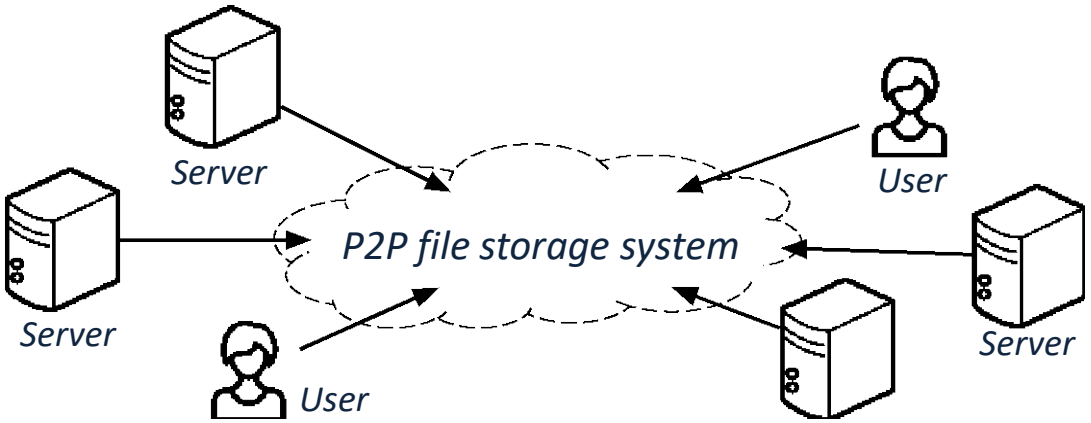

# Framework of blockchain-based P2P storage system

- P2P storage system

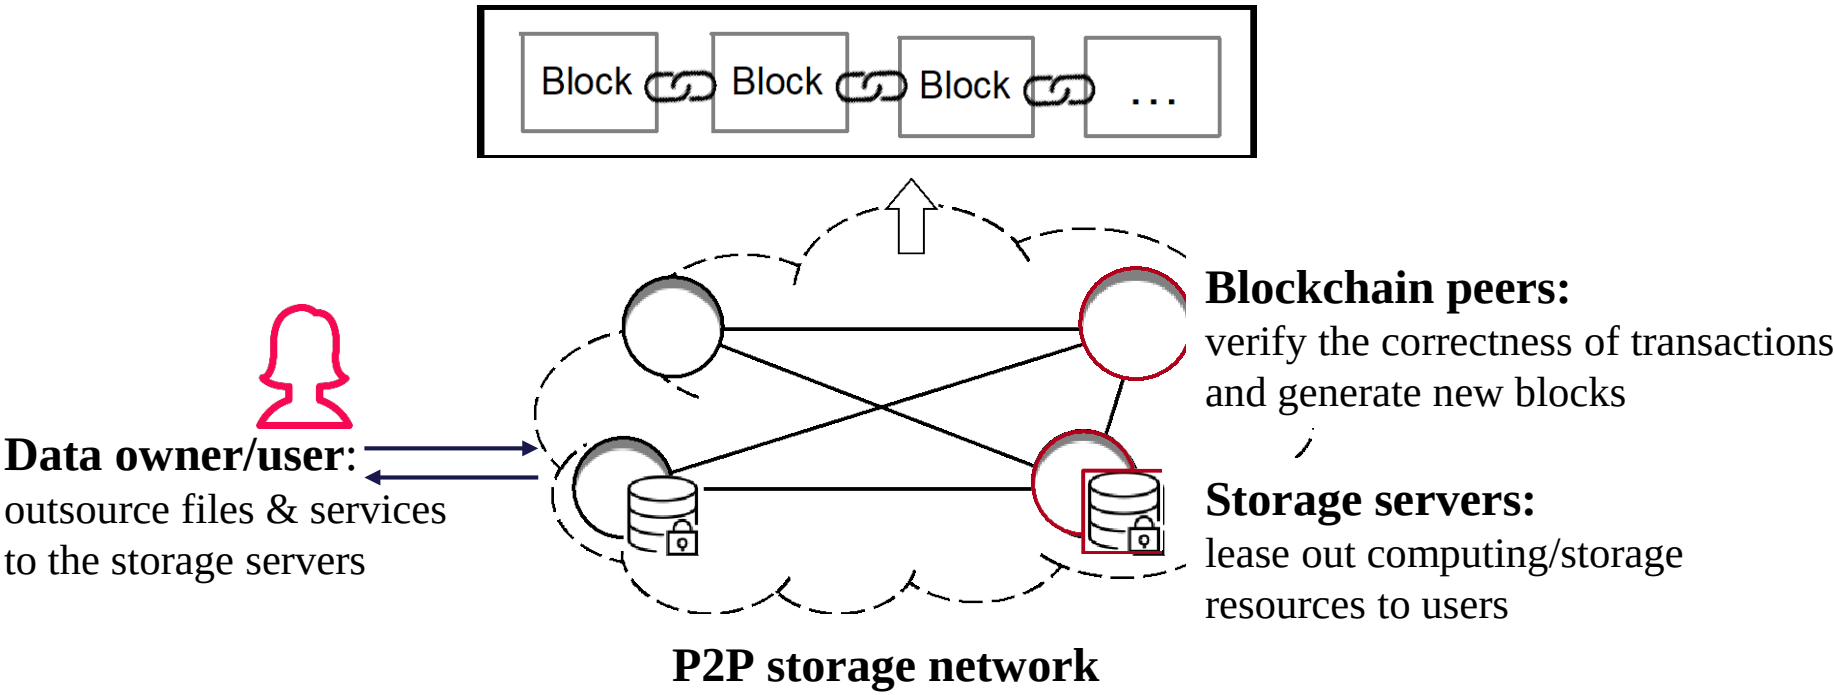


**A secure and fair platform for people to lease computing resources and for users to receive services**

# Framework of a blockchain-based P2P storage system

- Blockchain P2P network consists of storage servers and peers
  - storage servers can be peers
- Data owners/users interact with storage servers via transactions
  - data owners bind with servers via smart contracts
  - data and search indexes are stored off-chain at storage servers
  - all operations between owner/user and server are via transactions
    - contract transactions, data search/update transactions, etc
- Peers verify correctness of transactions and generate new blocks to the blockchain

# Signing a storage contract

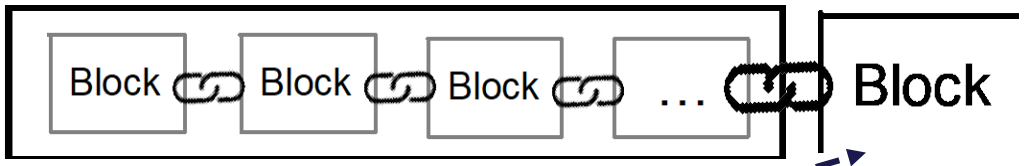


Storage Contract

- *Addr. of svr/owner*
- *Service fees*
- *Checklist*
- *Storage digest*

**Data owner**

- draft & sign the storage contract
- build & upload indexes and encrypted files to the server

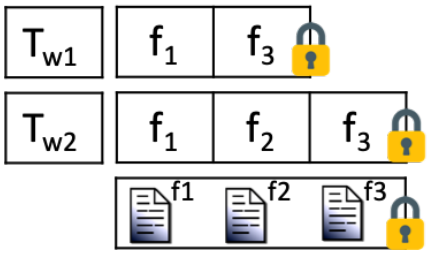


*contract trans.*
*contract details*

**Blockchain peers**

- verify the contract and record it in the chain

**Storage server**

- confirm & sign the storage contract
- broadcast the contract to blockchain network

# Search transaction

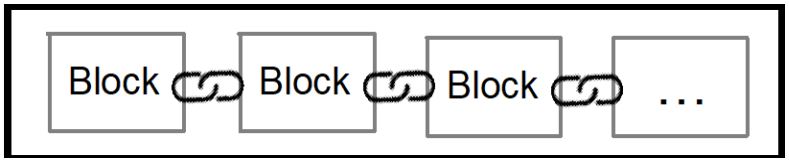


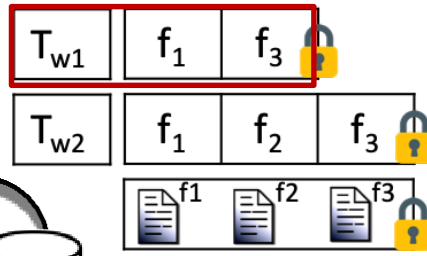


**User**

- generate a search trapdoor $\{T_{w1}\}$ for keyword $W_1$
- broadcast the trapdoor $\{T_{w1}\}$ to the blockchain network

*search trans.*
*trapdoor=*$\{T_{w1}\}$

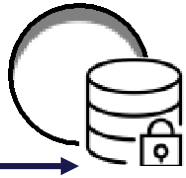

**Storage server**

- search the indexes
- broadcast results $\{f_1, f_3\}$ to the blockchain network

*search trans.*
*trapdoor=* $\{T_{w1}\}$

*result trans.*
*results=* $\{f_1, f_3\}$

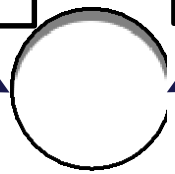

**Blockchain peers**

# Search result verification

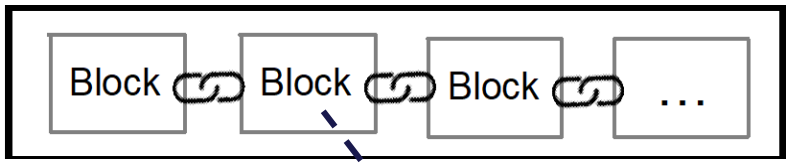



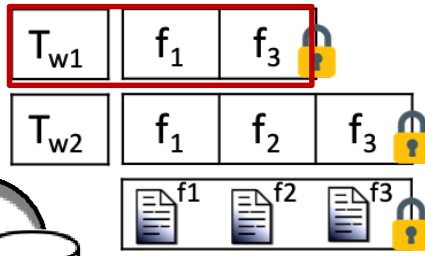


**User**

- broadcast the trapdoor $\{T_{w1}\}$ to the blockchain network

*search trans.*
*trapdoor= $\{T_{w1}\}$*

*result trans.*
*results= $\{f_1, f_3\}$*

**Storage server**

- broadcast results $\{f_1, f_3\}$ to the blockchain network

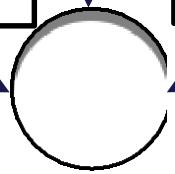

**Blockchain peers**

- verify the results against the on-chain checklist

# Generating new blocks to the blockchain

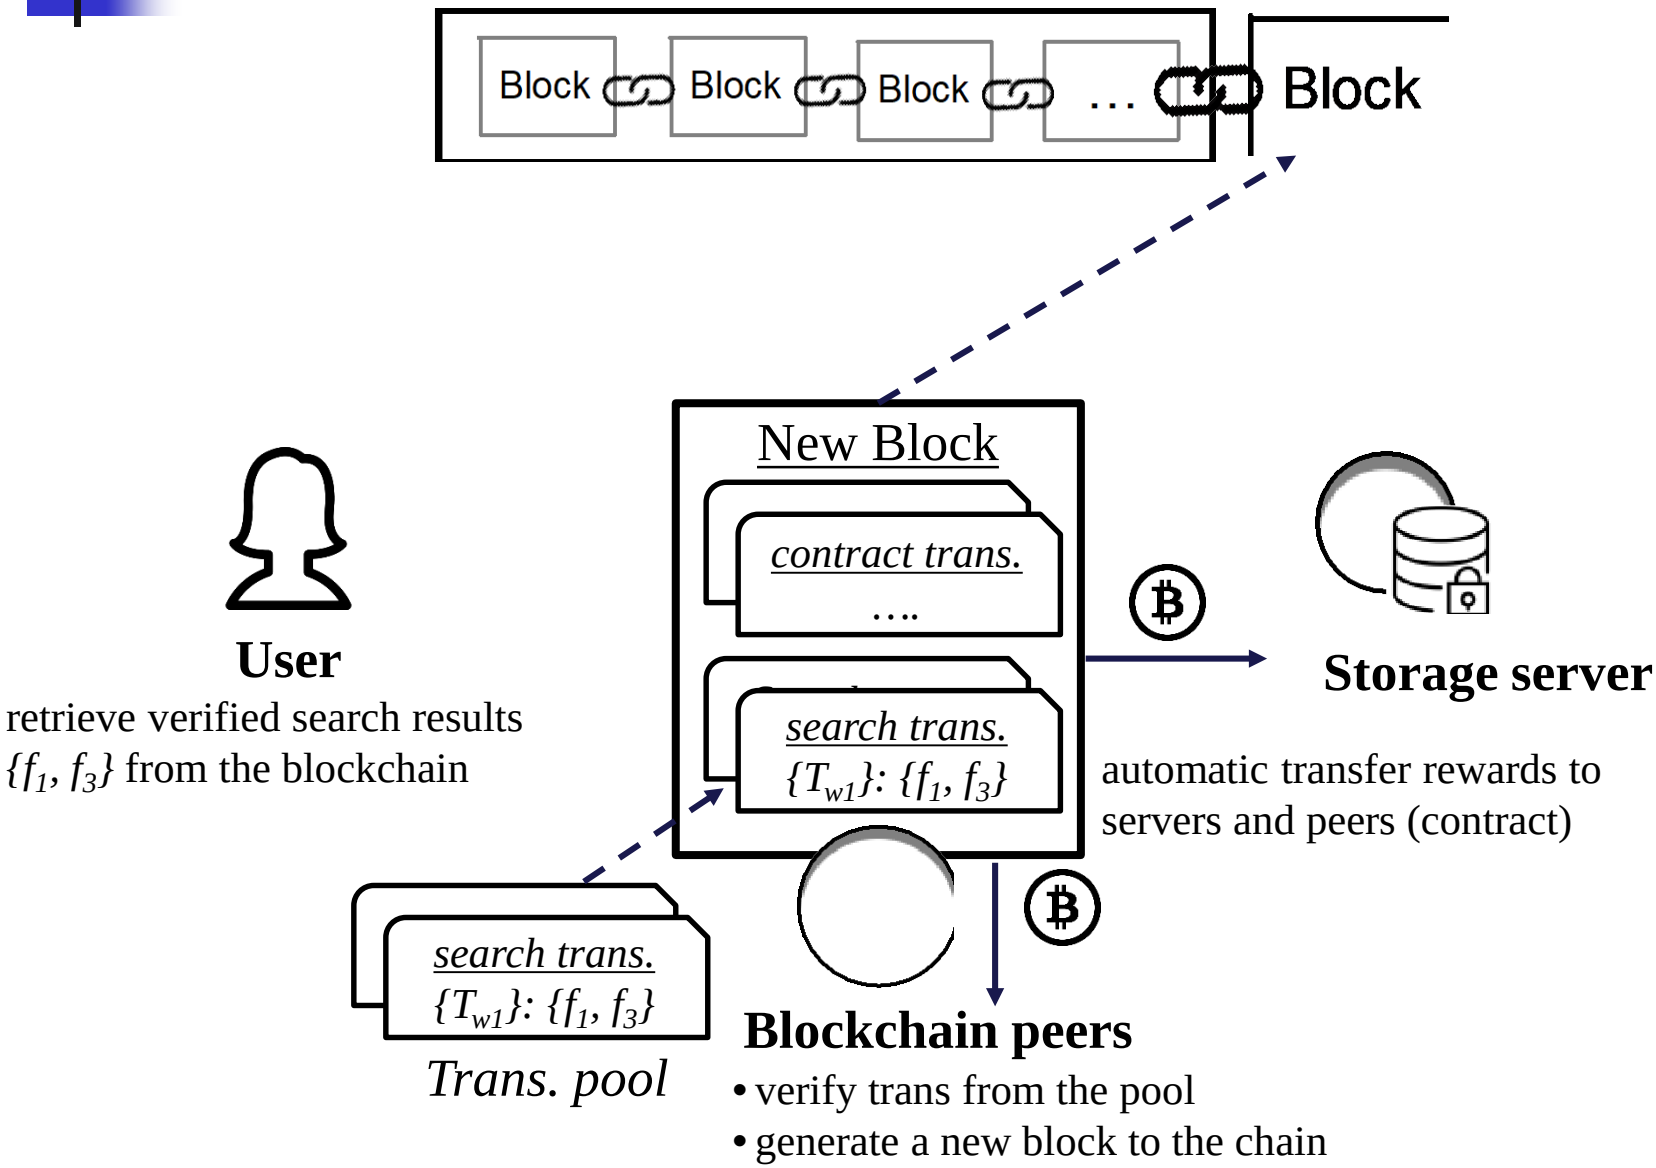

# A new consensus protocol

- Verification of a search result transaction includes:
  - verifying the search results, and
  - auditing the integrity of the stored file
- Peers compete with each other to generate new blocks

$$\text{SHA-256\^2}(\, pk \,\|\, T \,\|\, Mr(Tx) \,\|\, H(\pi) \,\|\, Block_{pre}\,) \overset{?}{<} Z \times B_{stc}$$

*peer ID* *ticket* *transitions* *file-proofs* *Prev. block* *peer's stake*

Mr(*Tx*): the Merkle-tree root of validated transactions in the new block

H($\pi$): the hash value of validated file-proofs

$B_{stc}$ : the peer's stake (amount of deposit it has in the system)

# A hybrid method of proof-of-stake and proof-of-work

- Proof-of-stake gives more advantage to peers with higher stake, reducing the average time for generating a new block
  - a trade-off between randomness and deterministic in block mining
  - increase the throughput of generating new blocks
- Peers perform data auditing as a useful PoW
- The longest chain rule still holds the global consensus among the peers